

# **CMPE110 Lecture 04**

# **Instruction Set Architecture**

---

Heiner Litz

<https://canvas.ucsc.edu/courses/19290>

# Announcements

---



- Homework groups have been created
- 1<sup>st</sup> HA Assignment will be released today
- Friday first Quiz


# Review

---





# Amdahl's Law



# Amdahl's Law: Make Common Case Efficient



- Given an optimization  $x$  that accelerates fraction  $f_x$  of program by a factor of  $S_x$ , how much is the overall speedup?

$$\text{Speedup} = \frac{\text{CPUTime}_{\text{old}}}{\text{CPUTime}_{\text{new}}} = \frac{\text{CPUTime}_{\text{old}}}{\text{CPUTime}_{\text{old}}[(1-f_x) + \frac{f_x}{S_x}]} = \frac{1}{(1-f_x) + \frac{f_x}{S_x}}$$

- Lesson's from Amdahl's law
  - Make common cases fast: as  $f_x \rightarrow 1$ , speedup  $\rightarrow S_x$
  - But don't overoptimize common case: as  $S_x \rightarrow \infty$ , speedup  $\rightarrow 1 / (1-f_x)$ 
    - Speedup is limited by the fraction of the code accelerated
    - Uncommon case will eventually become the common one
- Amdahl's law applies to cost, power consumption, energy ...



# Other Benchmark Suites

## ■ PARSEC

- Focus on Multicore/Multithreaded

## ■ STAMP

- Transactional Memory

## ■ Cloudsuite

- Datacenter applications

## ■ Microbenchmarks

- FIO (block I/O)
- Stream (memory bandwidth)
- LINPACK (HPC, embarrassingly parallel)



# Summarizing Benchmarks

Arithmetic mean

$$\frac{1}{n} \sum_{i=1}^n T_i$$

Use with times, not with rates

Represents total execution time

Harmonic mean

$$\frac{n}{\sum_{i=1}^n \frac{1}{R_i}}$$

Use with rates not with times

Geometric mean

$$\left( \prod_{i=1}^n \frac{T_i}{T_{ri}} \right)^{\frac{1}{n}} = \exp \left( \frac{1}{n} \sum_{i=1}^n \log \left( \frac{T_i}{T_{ri}} \right) \right)$$

Good with normalized performance

Does not represent total execution time

- Can also use weighted version
- Be careful which one you use!

# Example: SPECpower\_ssj2008 for Xeon X5650



## ■ Power/performance benchmark

$$\text{Overall ssj_ops per Watt} = \left( \sum_{i=0}^{10} \text{ssj\_ops}_i \right) / \left( \sum_{i=0}^{10} \text{power}_i \right)$$

Target Load %	Performance (ssj_ops)	Average Power (Watts)
100%	865,618	258
90%	786,688	242
80%	698,051	224
70%	607,826	204
60%	521,391	185
50%	436,757	170
40%	345,919	157
30%	262,071	146
20%	176,061	135
10%	86,784	121
0%	0	80
Overall Sum	4,787,166	1,922
$\Sigma \text{ssj\_ops} / \Sigma \text{power} =$		2,490



# Measuring Performance (Linux)

## ■ *time application*

- Measures execution time of the application
- Distinguishes between user and kernel

## ■ *perf record*

- Low overhead sampling
- Lists time spent per function

## ■ *perf topdown*

- Uses hardware performance counters (PMU)
- Enables microarchitectural studies



# Perf record/report


```
94.72%  perf-test  libm-2.19.so          [.] __sin_sse2
|
--- __sin_sse2
    |
    ---44.20%-- foo
        |
        ---100.00%-- main
                        __libc_start_main
                        _start
                        0x0

    --27.95%-- baz
        |
        ---51.78%-- bar
                        foo
                        main
                        __libc_start_main
                        _start
                        0x0

        --48.22%-- foo
                        main
                        __libc_start_main
                        _start
                        0x0
```



# Perf Topdown



Yasin, Ahmad. "A top-down method for performance analysis and counters architecture."



# Microarchitectural Simulation

- Simulators: Gem5, Zsim, Mars, HDL
- Emulation + Timing Model
- Models processors hardware
- Can probe everything -> detailed analysis
- Enables modeling of new techniques

# Instruction Set Architectures

---



# Before We Build Hardware

---



What is the HW/SW interface?

# At Their Core, Digital Systems Are Pretty Simple



Computers only work with binary signals (0 or 1 bits)

More complex stuff expressed as sequences of bits

Numbers, characters, strings, pictures, ...

Memory cells preserve bits over time

Flip-flops, registers, SRAM, DRAM


Logic gates operate on bits (AND, OR, NOT, multiplexor)

To get the HW to compute something

We express it as a sequence of simple instructions

We encode these instructions as strings of bits

# Big Picture: Running a Program



temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;  
  
ld \$15, 0(\$2)  
ld \$16, 4(\$2)  
sd \$16, 0(\$2)  
sd \$15, 4(\$2)

0000 1001 1100 0110 1010 1111 0101 1000  
1010 1111 0101 1000 0000 1001 1100 0110  
1100 0110 1010 1111 0101 1000 0000 1001  
0101 1000 0000 1001 1100 0110 1010 1111

High/Low on control lines




# Instruction Set Architecture (ISA)

- The HW/SW interface
  - The contract between HW and SW (compilers, assembler, etc)
- It defines
  - The state of the system (Registers, memory)
  - The functionality of each HW instruction
  - The encoding of each HW instruction
- It does ***not*** define
  - How instructions are implemented
  - How fast/slow instructions are
  - How much power instructions consume

# Instruction-Set Architecture

## Where does it fit?





# Does the ISA Choice Matter

- If we write software in high-level languages,  
ISA choice is irrelevant, right?
  - x86, ARM, ...
- Thoughts?



# Instruction Set Architecture (ISA)

Many different ISAs

Examples: Sun SPARC, PowerPC, IBM 390, MIP32, MIPS64, Intel x86 (IA32), Intel IA64 Intel x86-64, ARM A32, ARM A64

Many different chips can implement same ISA (family)

8086, 386, 486, Pentium, Pentium II, Pentium4 all implement IA32

ISAs last a long time, implementations are short lived

x86 has been in use since the 70s

IBM 390 started as IBM 360 in 60s

Stable interface for software (binary compatibility)

Can you change an ISA?

# RISC vs. CISC

## Complex/Reduced Instruction Set



- Hot debate in the 80'ies
- CISC: X86, IBM 360, Motorola 68K
- RISC: MIPS, SUN Sparc, RISC-V
- CISC:
  - $\text{MUL Mem2} \leq \text{Mem0} * \text{Mem1}$
- RISC:
  - Load REG0  $\leq$  Mem0
  - Load REG1  $\leq$  Mem1
  - $\text{MUL REG0} \leq \text{REG0} * \text{REG1}$
  - Store Mem2  $\leq$  Reg0



# RISC vs. CISC

CISC	RISC
Emphasis on Hardware	Emphasis on Software
Multi-cycle complex instructions	Simple (single clock) instructions
Memory-to-Memory load/store incorporated in instr.	Register-to-Register Separate load/store instructions
Small code size	Large code size
High CPI	Low CPI
Low clock frequency	High clock frequency
Variable length instructions	Same length instructions
Complex instruction decode	Simple instruction decode
HW difficult to implement	HW easy to implement



# Intel/AMD's take on CISC

- X86 is a CISC ISA
- Internally Intel/AMD CPUs are RISC
- Transform CISC instructions into uOps
- Can use microcode engine
- Small code footprint of CISC, simple RISC instrs
- But: Decoding is still an issue (variable length)



# Instruction Length and Format

## ■ Fixed Length

- Address of next instruction is easy to compute
- ***Code density***: common instructions as long as rare

## ■ Variable Length

- Better code density
  - x86 averages 3 bytes (from 1 to 16 per instruction)
  - Common instructions are shorter
  - Less instruction memory to fetch
- Fetch and decode are more complex
- Compromise: N fixed sizes (e.g. 32,64,128 bit)

# Working Example: RISC-V ISA

---





# RISC-V ISA

Developed by UC Berkeley

Fully Open-source

Clean slate design

Typical of RISC ISAs (e.g., ARM, SPARC, MIPS)

Combines the best of prior RISC ISAs

Google, Nvidia, Qualcomm, Samsung, NXP, Micron, Marvell

Why RISC-V instead of Intel x86?

RISC-V is simple, elegant and easy to understand

Is becoming the 3<sup>rd</sup> most relevant ISA (after x86 & ARM)

x86 is ugly and complicated to explain

x86 is dominant on desktop

ARM > x86 but still ugly and complicated

# ISA Design Principles



Simplicity favors regularity



Smaller is faster

Good design demands good compromises

Make the common case fast



# RISC-V: System State






# RISC-V Assembly

## ■ Human readable form

### ■ Example: add x3, x2, x1



Add value of register x2 and x1 and store into x3

# RISC-V Instruction Format (machine format)



## R-type



Opcode: basic operation of instruction (7)

Rs1: Register source 1 operand (5)

Rd: Register destination operand (5)

Rs2: register source 2 operand (5)

Funct3: additional opcode field (3)

Funct7: additional opcode field (7)

**Question: Why did RISC-V only define 32 registers?**

I-type      **Question: What is an immediate?**



S/B-type      **Question: Why is the immediate split?**



U/J-type






# RISC-V Instruction Format

- All instructions are 32 bit to alleviate decoding
  - Smaller is faster
- Requires to interpret bit fields differently for different instructions (R vs. I vs. S/B vs. U/J)
  - Simplicity favors regularity
- Limits register file size to 32 (5 bits per operand)
  - Good designs demand good compromises




# Instruction Execution



Before State



# Instruction Execution




Before State



# Instruction Execution

Add x3, x1, x2

31	25 24	20 19	15 14	12 11	7 6	0
0000000	00001	00010	000	00011	0110011	




Before State




# Instruction Execution

Add x3, x1, x2

31	25 24	20 19	15 14	12 11	7 6	0
0000000	00001	00010	000	00011	0110011	



Before State



After State



# Instruction Execution

Start with the *before state* of the machine

PC, Regs, Memory

PC used to *fetch* an *instruction* from memory

Instruction directs how to compute the *after state* of the machine

For add:  $PC = PC + 4$  and  $rd = rs + rt$

This happens *atomically*

# C vs. RISC-V Programmers Interface



	<b>C</b>	<b>RISC-V ISA</b>
Registers		32 64b integer, R0 = 0 32 32b single FP 16 64b double FP PC and special registers
Memory	local variables global variables	$2^{61}$ linear array of bytes
Data types	int, short, char, unsigned, float, double, aggregate data types, pointers	doubleword(64b), word(32b), byte(8b), half-word(16b)
Arithmetic operators	+ , - , * , % , ++ , < , etc.	add, sub, and, sll, etc.
Memory access	a, *a, a[i], a[l][j]	ld, sd, lh, sh, lb, sb
Control	If-else, while, do-while, for, switch, procedure call, return	branches, jumps, jump and link



# Why Have Registers?

## Alternative: memory-memory ISA?

All HLL(C/Java/..) variables declared in memory

Instructions operate directly on memory operands?

E.g. Digital Equipment Corp (DEC) VAX ISA

## Benefits of registers

Smaller is faster (100-1000x)


Multiple concurrent accesses

Shorter names (fewer bits to encode)

## Load-Store or RISC ISAs

Arithmetic instructions only use register operands

Data loaded into registers, operated on, and stored back to memory





# Using Registers

Registers are a finite resource that needs to be managed

By the assembling programmer

Or the compiler (register allocation)

## Goal

Keep data in registers as much as possible

## Issues

Finite number of registers available

Spill registers to memory when all registers in use

## Arrays

Data is too large to store in registers

What's the impact of fewer or more registers?



# Working Code Example (C)

Simple C procedure: sum\_pow2 =  $2^{b+c}$

```
1: int sum_pow2(int b, int c)
2: {
3:     int pow2 [8] = {1, 2, 4, 8, 16, 32, 64, 128};
4:     int a, ret;
5:     a = b + c;
6:     if (a < 8)
7:         ret = pow2[a];
8:     else
9:         ret = 0;
10:    return(ret);
11: }
```



# Arithmetic Instructions

Consider C statement:  $a = b + c$ ; (line 5)

Assume the variables are in registers  $x_1-x_3$  respectively

**Add x1 , x2 , x3                          # a = b + c**

Similar instructions

Arithmetic: sub

Logical: and, or, nor



# Complex Operations

What about more complex C statements?

```
a = b + c + d - e;
```

Break into multiple instructions

```
Add xt0, x1, x2      # x5 = b + c
```

```
Add xt0, xt0, x3      # x5 = x5 + d
```

```
Sub xt0, xt0, x4      # a = x5 - e
```

xt0 is a temporary register

# Numbers representation: Signed & Unsigned



If given  $b[n-1:0]$  in a register or in memory

Unsigned value

$$value = \sum_{i=0}^{n-1} b_i 2^i$$

Signed value (2's complement)

$$value = -(b_{n-1} 2^{n-1}) + \sum_{i=0}^{n-2} b_i 2^i$$



# Unsigned & Signed Numbers

X	unsigned	signed
0000	0	0
0001	1	1
0010	2	2
0011	3	3
0100	4	4
0101	5	5
0110	6	6
0111	7	7
1000	8	-8
1001	9	-7
1010	10	-6
1011	11	-5
1100	12	-4
1101	13	-3
1110	14	-2
1111	15	-1

## Example values

4 bits

Unsigned:  $[0, 2^4 - 1]$

Signed:  $[-2^3, 2^3 - 1]$

## Equivalence

Same encodings for non-negative values

## Uniqueness

Every pattern represents unique integer

Not true with sign magnitude



# RISC-V Constants

Often want to specify a constant operand in the instruction

Constant == immediate == literal == offset

Use the `addi` instruction

**addi dst, src1, immediate**

The immediate is a 12 bit signed value between  $-2^{11}$  and  $2^{11}-1$

To enable 32 bit immediate use **lui dst, immediate**

Example:

C: `a++;`

RISC-V: `addi x1, x1, 1 # a = a + 1`

# Keep in Mind: Arithmetic Overflow




When the sum of two n-bit numbers can not be represented in n bits

## Unsigned


- Wraps Around
  - If true sum  $\geq 2^n$
  - At most once

True Sum



## Signed

True Sum





# Memory Data Transfer


Data transfer instructions move data to and from memory

Load moves data from a memory location to a register

Store moves data from a register to a memory location

All memory access happens through loads and stores

Floating-point loads and stores for accessing FP registers





# Data Transfer Instructions: Loads

Data transfer instructions have three parts

Operator name (defines the transfer size as well)

Destination register


Base register address and constant offset

`ld dst, offset(base)`

Offset value is a 12-bit signed constant (immediate)



# Displacement Addressing Mode



Effective address is a byte addresses

Must be aligned to words, half-words, & bytes

More on this later



# Loading Data Example

Consider the C example:  $a = b + *c;$

Assume a in x1, b in x2, c in x3

ld instruction:

```
Ld xt0, 0(x3)      # xt0 = Memory[c]  
                      # xt0 is temp reg
```

```
Add x1, x2, xt0 # a = b + *c
```



# Accessing Arrays

Arrays are really pointers to the base address in memory

Address of element A[0]

Use offset value to indicate which index

Note: addresses are in bytes, so multiply by the size of the element

Unlike C, assembly does not do pointer arithmetic for you!

Consider the integer array where `pow2` is the base address


With this compiler on this architecture, each int requires 4 bytes

The data to be accessed is at index 5: `pow2 [5]`

Then the address from memory is `pow2 + 5 * 4`



# Array Memory Diagram





# Array Example

Example: `a = b + pow2[7];`

`x3 = 1000`

`lw` instruction with offset:

```
Ld xt0, 28(x3)    # xt0 = Memory[pow2[7]]  
Add x2, x1, xt0  # a = b + pow2[7]
```



# Storing Data

Store is the reverse of load

And identical in address generation

Copy data from the source register to an address in memory

**sd src, offset(base)**

Offset value is a 12-bit signed constant



# Storing Data Example

Example: **\*a = b + c;**

Assume a in x3, b in x1, c in x2

sd instruction:

<b>add xt0, x1, x2</b>	<b># \$t0 = b + c</b>
<b>sd xt0, 0(x3)</b>	<b># Memory[s0] = b + c</b>



# Storing to an Array

Example: **a[3] = b + c;**

Assume b in x2, c in x3, a in x4

sd instruction with offset

```
add xt0, x2, x3      # $t0 = b + c  
sd xt0, 12(x4)      # Memory[a[3]] = b + c
```

Question: Why 12 and not 3?



# Indexed Array Storage

Example:  $a[i] = b + c;$

Assume i in x3, b in x1, c in x2, a in x4

Address generation + store:

```
Add xt0, x1, x2# $t0 = b + c
Sll xt1, x3, 2          # $t1 = 4 * i
Add xt2, x4, xt1        # $t2 = a + 4*i
sd  xt0, 0(xt2) # Memory[a[i]] = b+c
```



# Interface to I/O

Load/stores provide interface to memories

How about interface to I/O devices?

Huge variety in I/O devices

Printer, USB camera, network interface, hard disk....

Huge variety in functionality and performance requirements

Special I/O instructions for each type of device

What would be the problem with this?



# Interface to I/O

Break the problem in two

- Communicating bits to the I/O device

- Executing operations based on these bits

## Register/memory based interface to I/O

- Every device includes some registers and memory

- Reading/writing these registers communicates bits

  - Similar across all I/O devices

- Every device has its own protocol of what these bits mean

  - Check status, read command, send command, ...

  - Software must know this protocol (device driver)




# Memory Mapped I/O: Preview

Loads/stores can be used to move bits to/from I/O devices

A load moves data from a I/O device register to a CPU register

A store moves data from a CPU register to a I/O device register



I/O register at address 0x80000000