

CMPE110 Lecture 13

Pipelining 3

Heiner Litz

<https://canvas.ucsc.edu/courses/19290>

Announcements



- Evaluation of the Tensor Processing Unit: A Deep Neural Network Accelerator for the Datacenter
- David Patterson, UC Berkeley & Google
- Friday, February 8, 2:40 PM, E2-180

Announcements



- No Quiz on Friday
- Midterm: Wed 02/13
 - 1 double sided page of notes
 - RISC-V Green card/cheat sheet part of the exam

Review



Forwarding Paths: Example



Example 1:

add \$1,\$1,\$2

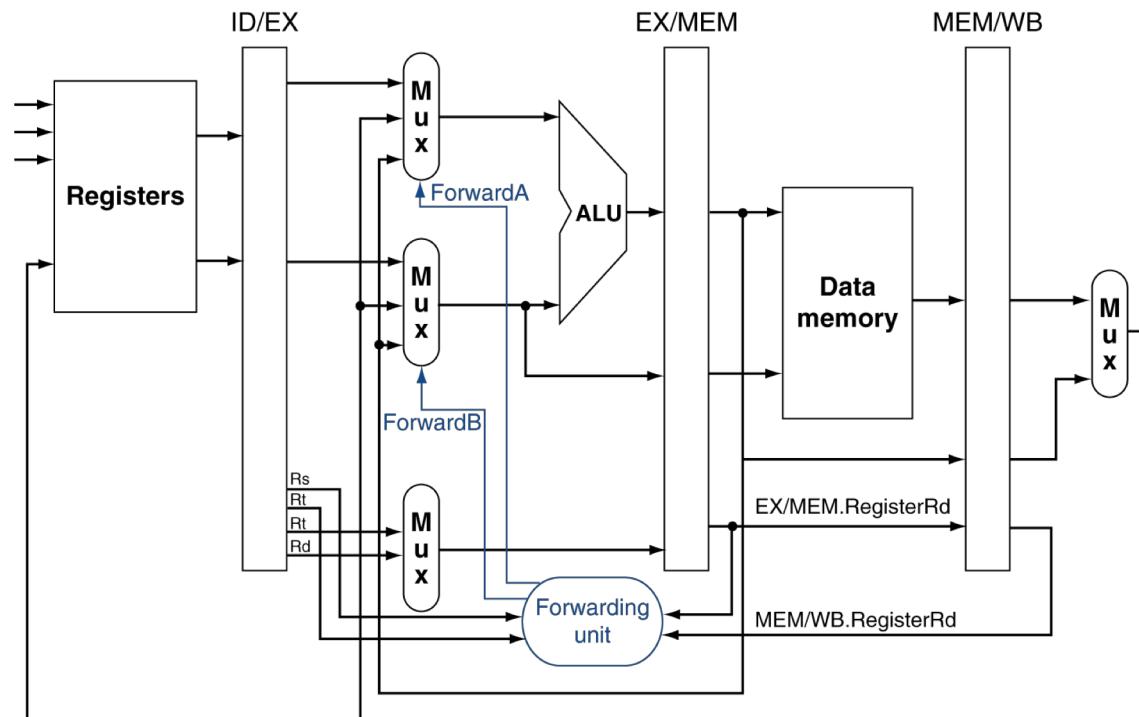
or \$1,\$1,\$4

Example 2:

add \$1,\$1,\$2

sub \$4,\$5,\$3

or \$1,\$1,\$4



b. With forwarding

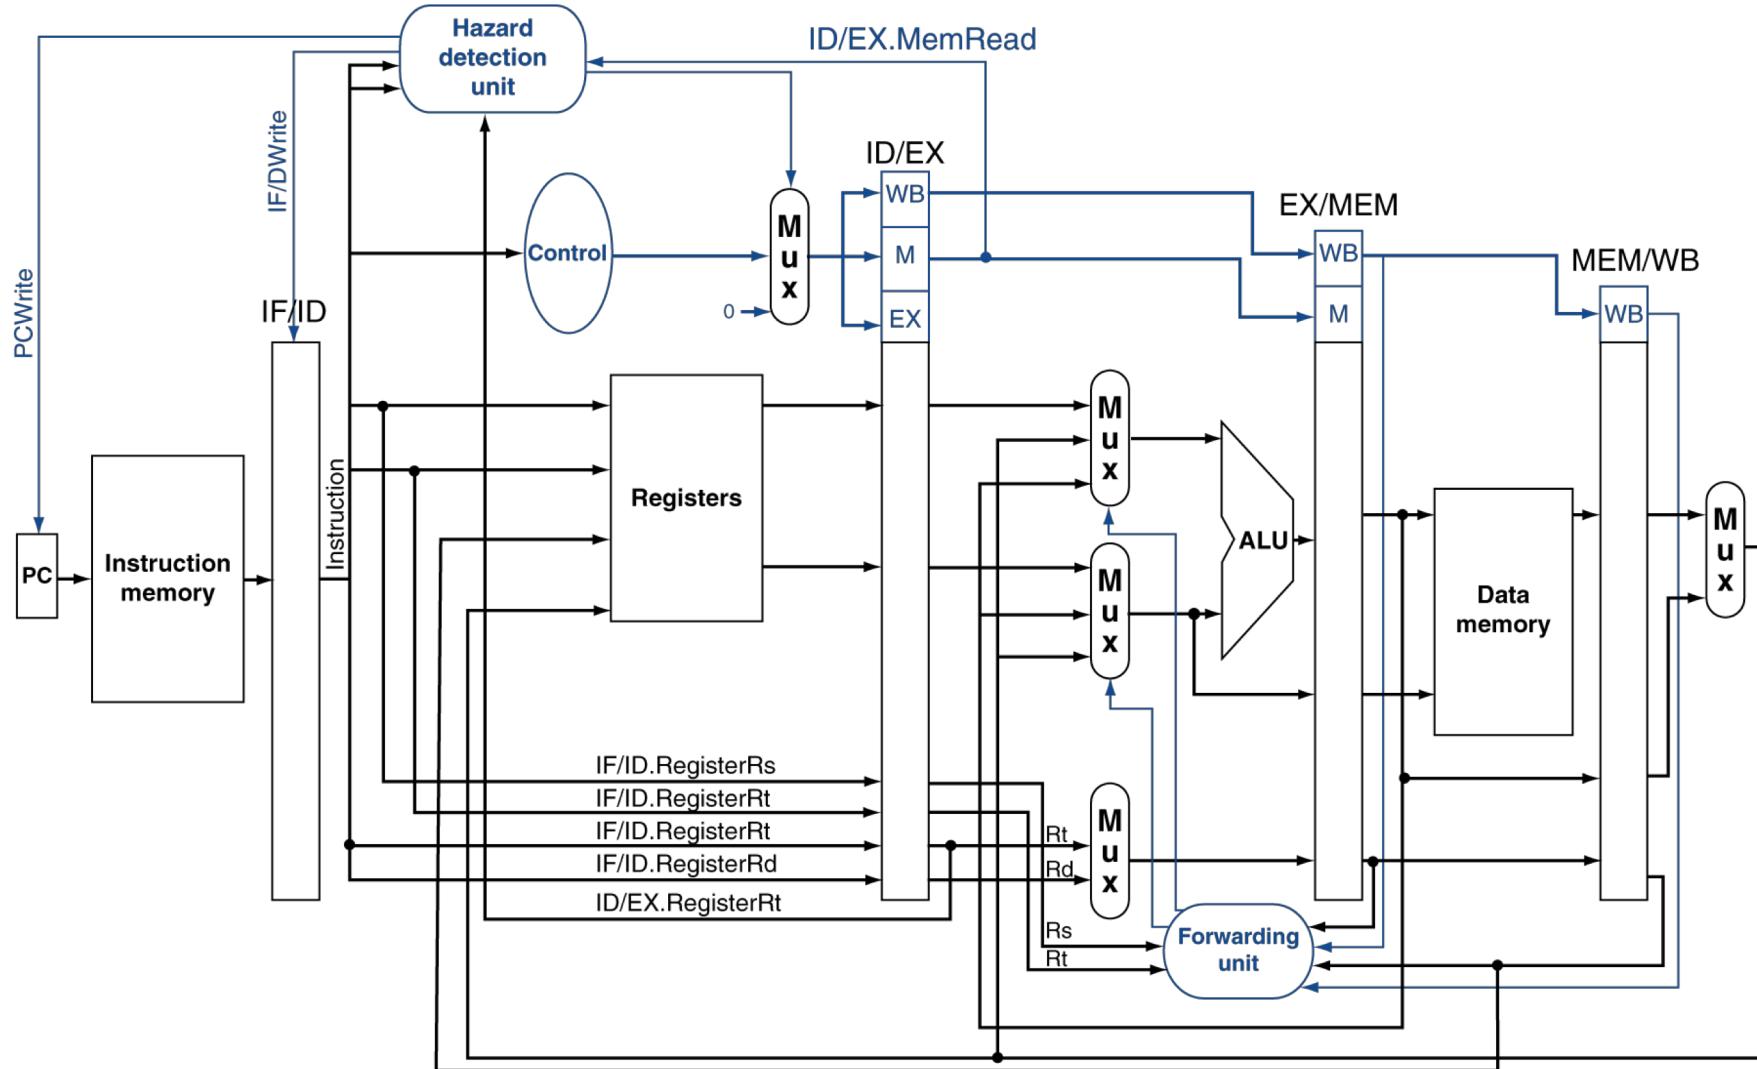
Example 3:

ld \$1,0(\$2)

sub \$4,\$5,\$3

or \$1,\$1,\$4

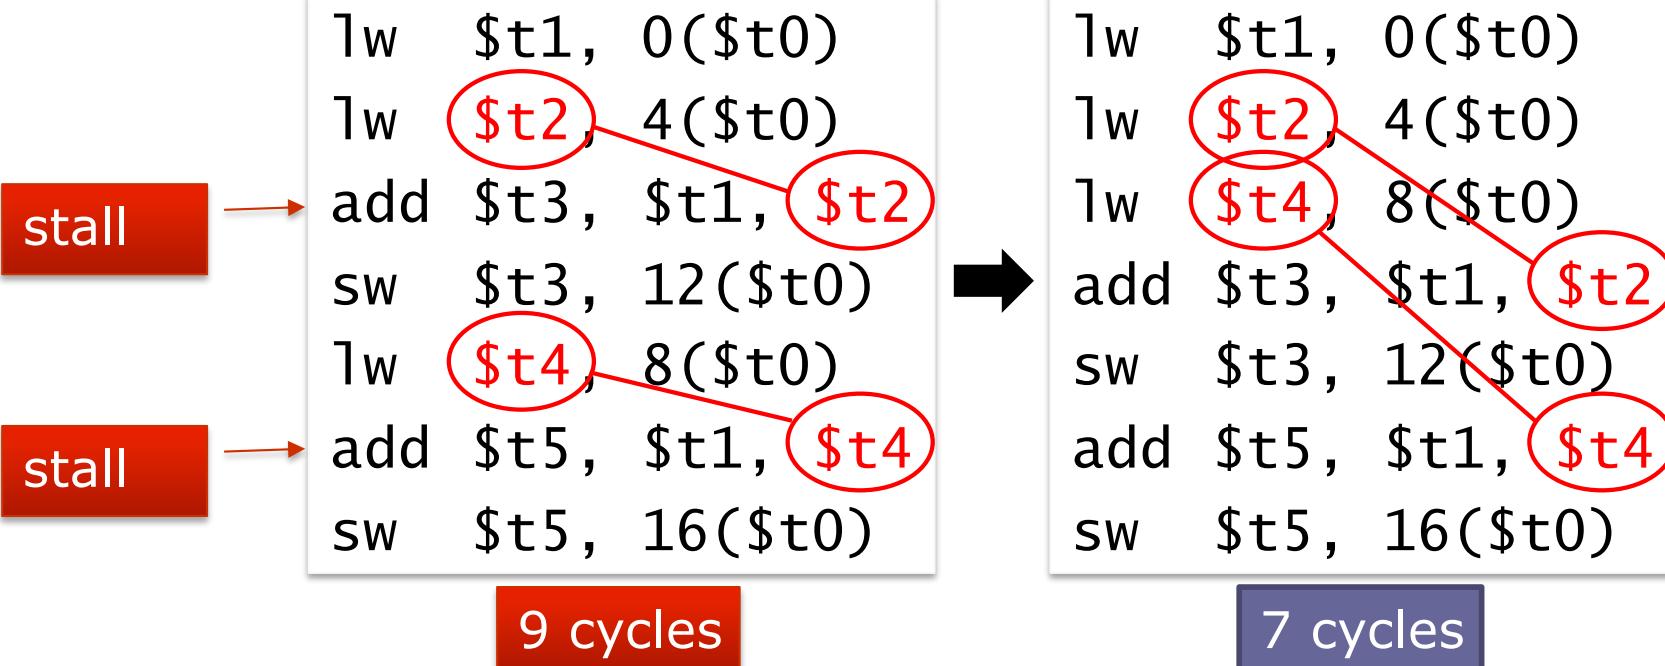
Datapath with Hazard Detection



Code Scheduling to Avoid Stalls



- Reorder code to avoid use of load result in the next instruction
- C code for $A = B + E; C = B + F;$



Control Hazards



- Branches get computed in the EX stage
- Requires 2 NOPs after every branch ☹
- Can we do better?



Reducing Branch Delay

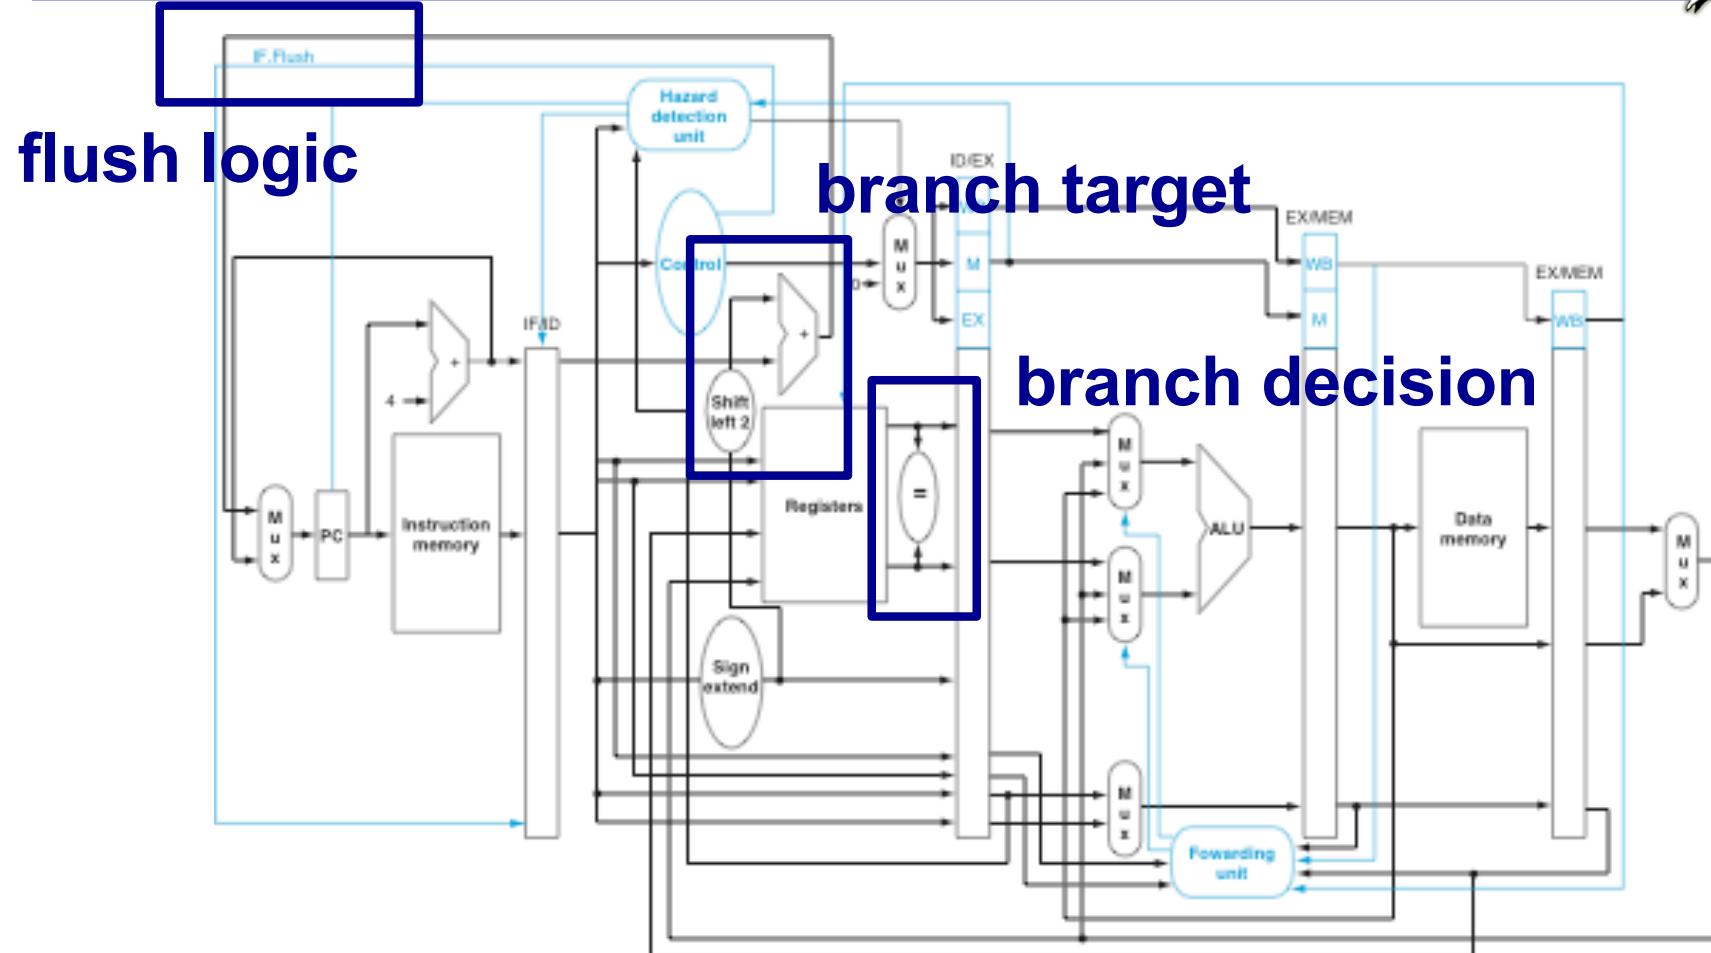
■ Minimize “bubble” slots

- Move branch computation earlier in the pipeline
- branch outcome: add comparator to ID stage
- branch target: add adder to ID stage

■ Predict branch not taken

- if correct, no bubbles inserted
- if wrong, flush pipe, inserting one bubble

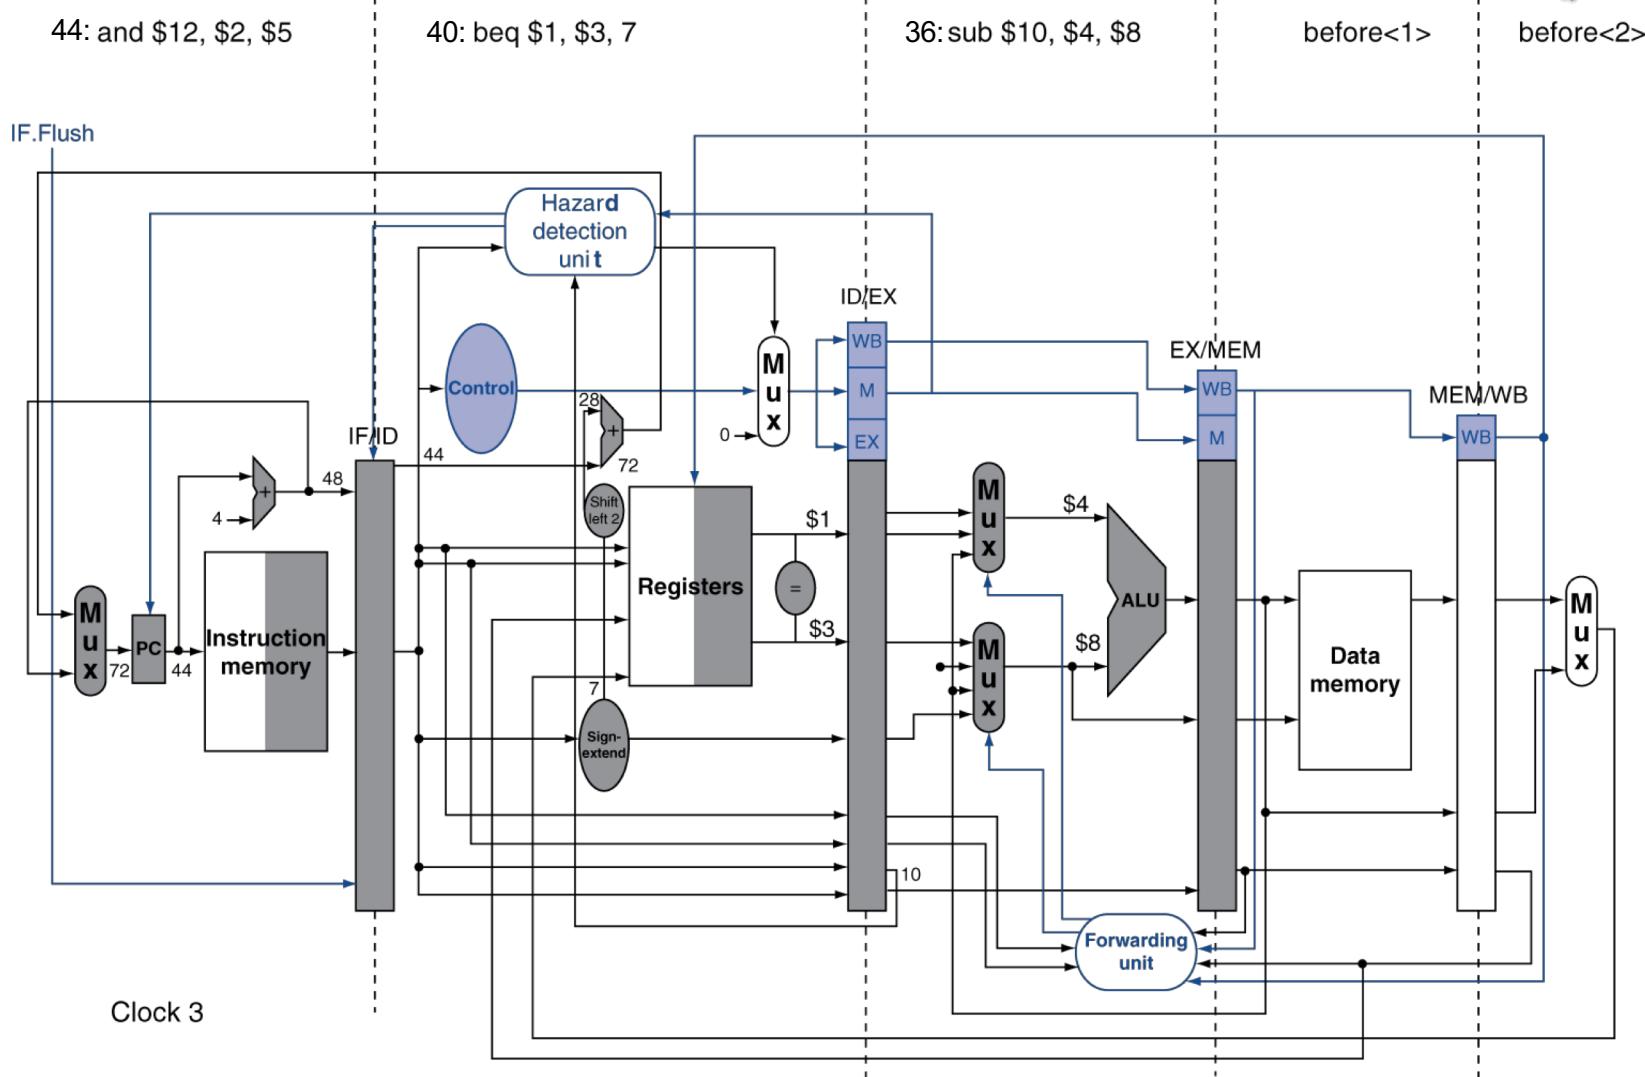
Reducing Branch Delay



move branch computation earlier in the pipeline
enable flush capability if mispredicted

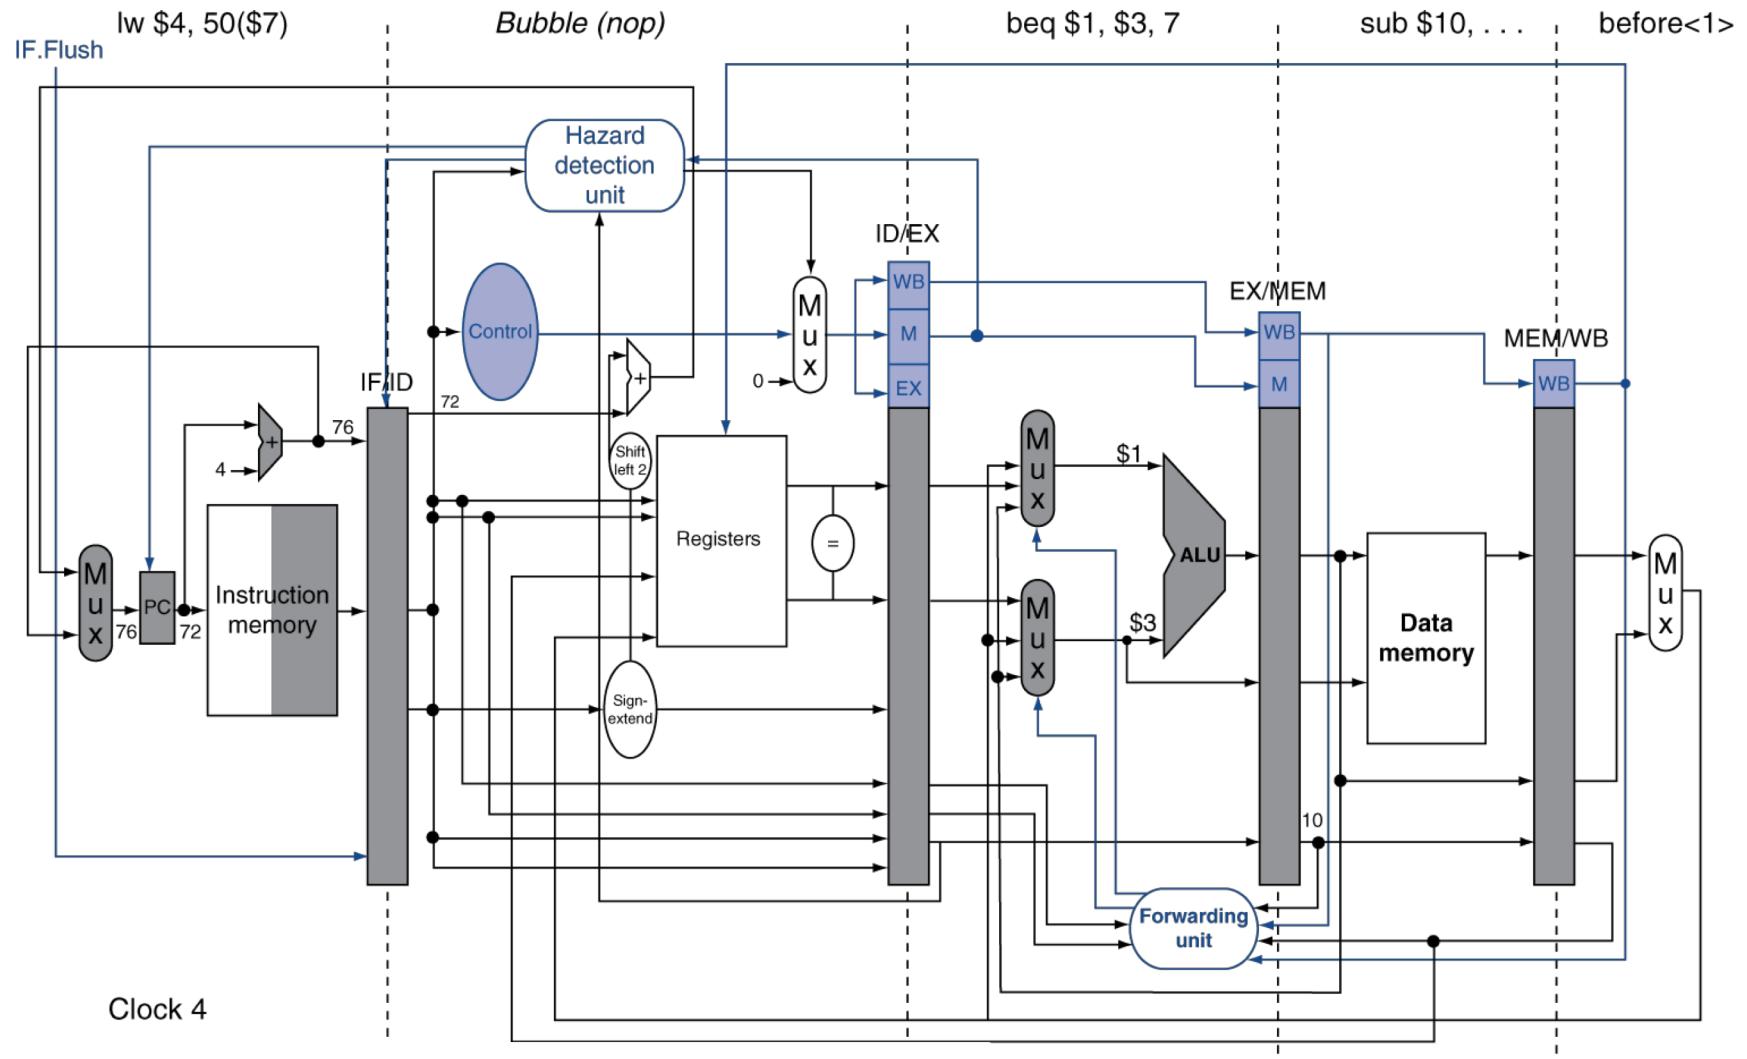


Example: Branch Taken





Example: Branch Taken



`nop inserted using IF.Flush`



How to Flush?

■ 5-stage Pipeline

- only one instruction to flush
- Can do by “flushing” the output register of IF

■ Deeper Pipelines

- Multiple cycles to determine branch outcome
- Need to know how many instructions in the pipe
- Flush all state changing actions
 - RegWrite, MemWrite, PCWrite (jmp, beq), overflow..

Performance Impact of Branch Stalls



- Need to stall for one cycle on every branch
- Consider the following case
 - The ideal CPI of the machine is 1
 - The branch causes a stall
- Effective CPI if 15% of the instructions are branches?
 - The new effective CPI is $1 + 1 \times 0.15 = 1.15$
 - The old effective CPI was $1 + 3 \times 0.15 = 1.45$



Delayed Branches

- An ISA-based solution used in early MIPS machines
 - Had a branch delay of one

- Example:

beq r1, r2, L

Branch instruction

sub r4, r1, r3

This operation ALWAYS is executed

and r6, r2, r7

This operation executes if branch fails

- Worked well initially

- Compiler can fill one slot 50% of the time (50% nops)

- For modern processors, it is a pain

- Many delay slots needed due to deeper pipelines
 - Processors deal with this using branch prediction
 - Delayed branches complicate exceptions as well

It's All Connected

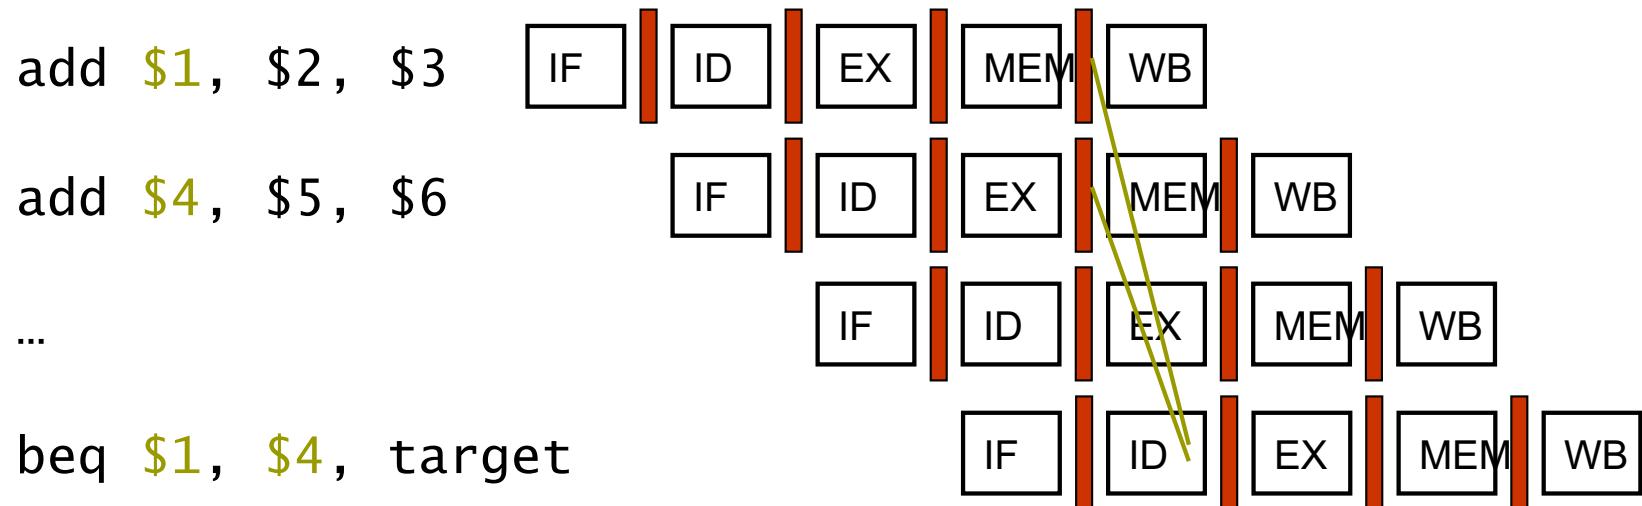


- We moved the branch control point to ID
- Which means we consume two registers in ID
 - earlier in the pipeline as before (EX)
- What does this mean for forwarding/stalls?
 - Need to check!!



Data Hazards for Branches

- If a comparison register is a destination of 2nd or 3rd preceding ALU instruction

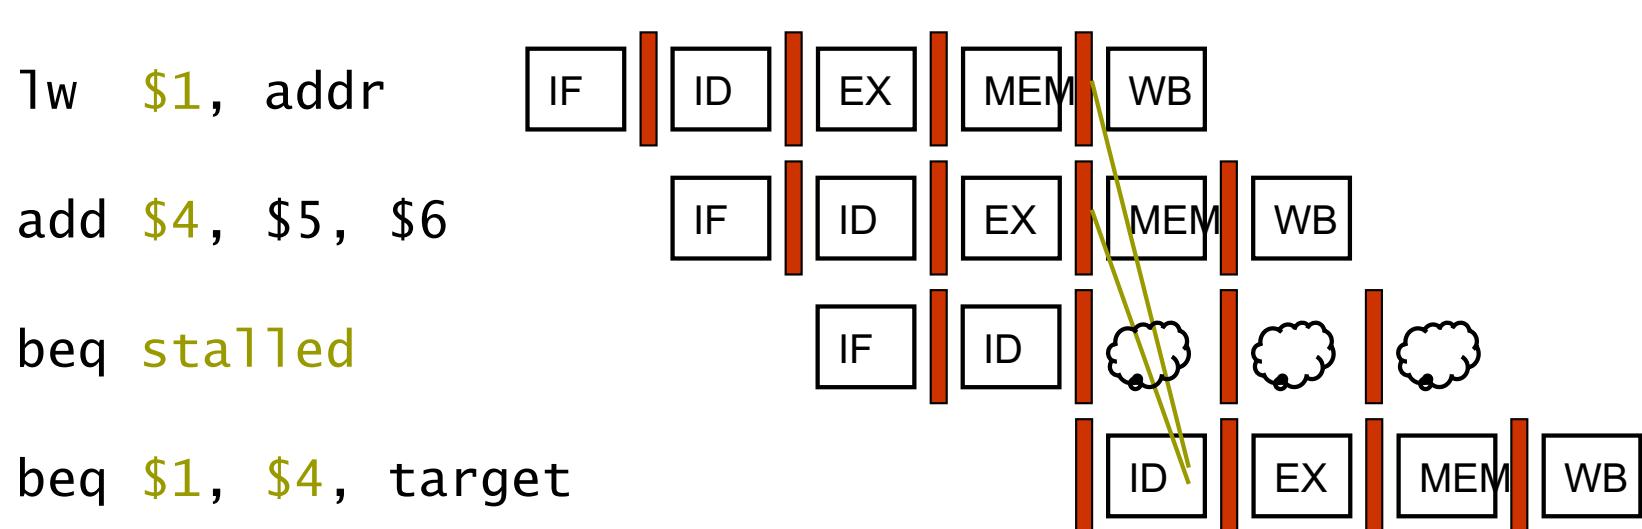


- Can resolve using forwarding
 - Additional datapaths and control



Data Hazards for Branches

- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle





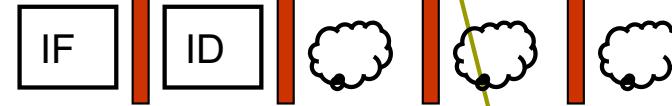
Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles

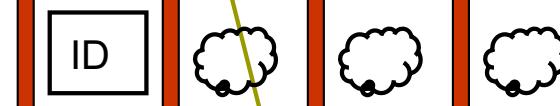
lw \$1, addr



beq stalled



beq stalled



beq \$1, \$0, target



Branch CPI



15% Branch frequency

65% branches are taken

50% of single delay slots are filled usefully

Control point	Branch strategy	Branch Stall CPI
MEM	stall	
ID	stall	
ID	Predict taken	
ID	Predict Not taken	
ID (target) / MEM (==)	Predict taken	
ID	Delayed branch	



Branch CPI

15% Branch frequency

65% branches are taken

50% of single delay slots are filled usefully

Control point	Branch strategy	Branch Stall CPI
MEM	stall	$3 \cdot 0.15 = 0.45$
ID	stall	$1 \cdot 0.15 = 0.15$
ID	Predict taken	$1 \cdot 0.15 = 0.15$
ID	Predict Not taken	$(1 \cdot 0.15) \cdot 0.65 + 0 \cdot 0.35 = 0.098$
ID (target) / MEM (==)	Predict taken	$(1 \cdot 0.15) \cdot 0.65 + (3 \cdot 0.15) \cdot 0.35 = 0.255$
ID	Delayed branch	$(1 \cdot 0.15) \cdot 0.50 = 0.075$

Branch Prediction



■ Static prediction vs. dynamic prediction

■ Static prediction schemes:

- always predict taken
- always predict not-taken
- compiler/programmer hint
- if ($\text{target} < \text{PC}$)
 - predict taken
 - else
 - predict not-taken

What's the rationale
behind this?



Dynamic Branch Prediction

■ Branch History Table (BHT)

- One entry for each branch PC
- Taken/Not taken bit

■ Branch Target Buffer (BTB)

- One entry for each branch PC
- Target address

■ Increasingly important for long pipelines (ID_x)

- x86 vs. RISC-V instruction decode



Branch Prediction Example

```
void foo() {  
    for(e=0;e<4;e++) {  
        stuff ...  
    }  
}  
  
foo:j E  
L:  stuff..  
    ..  
E:  bne $t3,$t4,L
```

How many times will bne at E: be predicted correctly?

Two-bit saturating counter:
strongly taken, weakly taken, weakly n-t, strongly n-t