

CMPE110 Lecture 14

Exceptions

Heiner Litz

<https://canvas.ucsc.edu/courses/19290>

Announcements

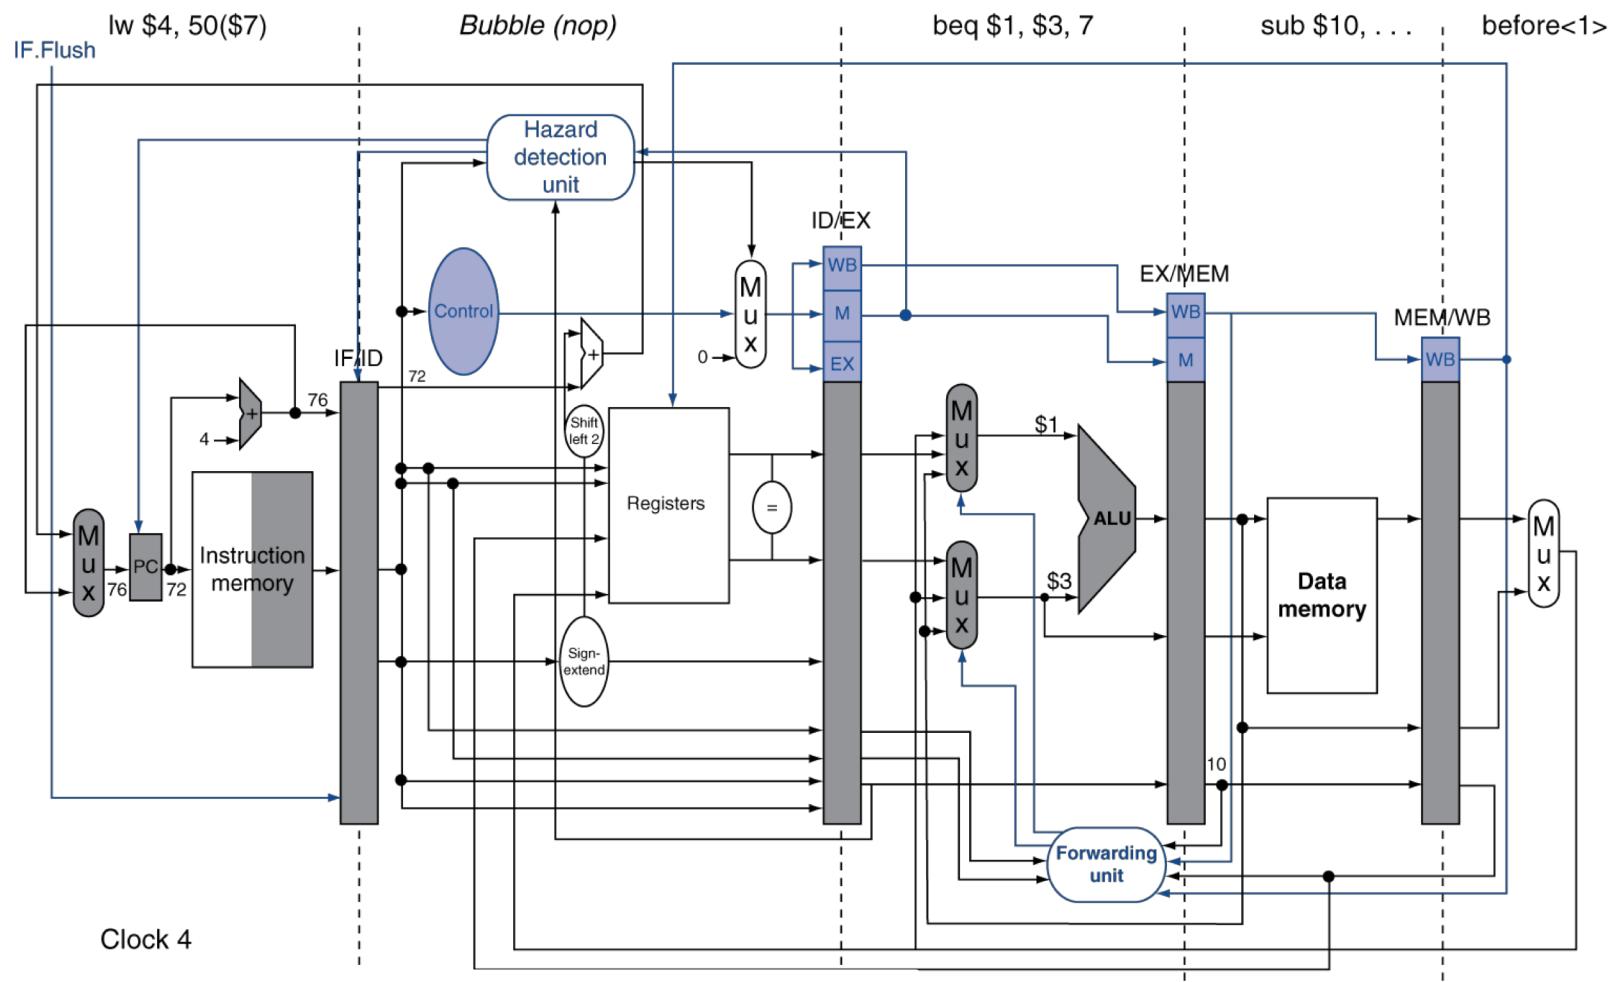


- No Quiz on Friday
- Midterm: Wed 02/13
 - 1 double sided page of notes
 - RISC-V Green card/cheat sheet part of the exam

Review



Example: Branch Taken



nop inserted using IF.Flush

Data Hazards for Branches

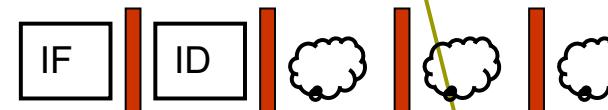


- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles

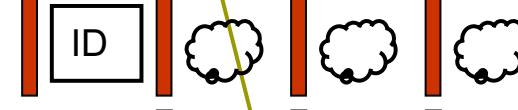
lw \$1, addr



beq stalled



beq stalled



beq \$1, \$0, target



Exceptions and Interrupts



Exceptions and Interrupts



- “Unexpected” events requiring change in flow of control
 - Switch from user to privileged (kernel) mode
 - Different ISAs use the terms differently
- Exception
 - Arises within the CPU
 - E.g., undefined opcode, overflow, div by 0, syscall, ...
- Interrupt
 - E.g., from an external I/O controller (network card)
- Dealing with them without sacrificing performance is hard

Handler Actions



- Read cause, and transfer to relevant handler
 - A software switch statement
 - May be necessary even if vectored interrupts are available
- Determine action required
- If user process is restartable
 - Take corrective action
 - Use EPC to return to program
- Otherwise
 - Terminate program (e.g., segfault, ...)
 - Report error using EPC, cause, ...

Precise Exceptions



- Definition: precise exceptions
 - All previous instructions had completed
 - The faulting instruction was not started
 - None of the following instructions were started
 - No changes to the architecture state (registers, memory)
- Why are they desirable by OS developers?
- With single cycle design, precise exceptions are easy
 - Why?
- With pipelined design, they can be tricky
 - Why?

Exceptions in a Pipeline



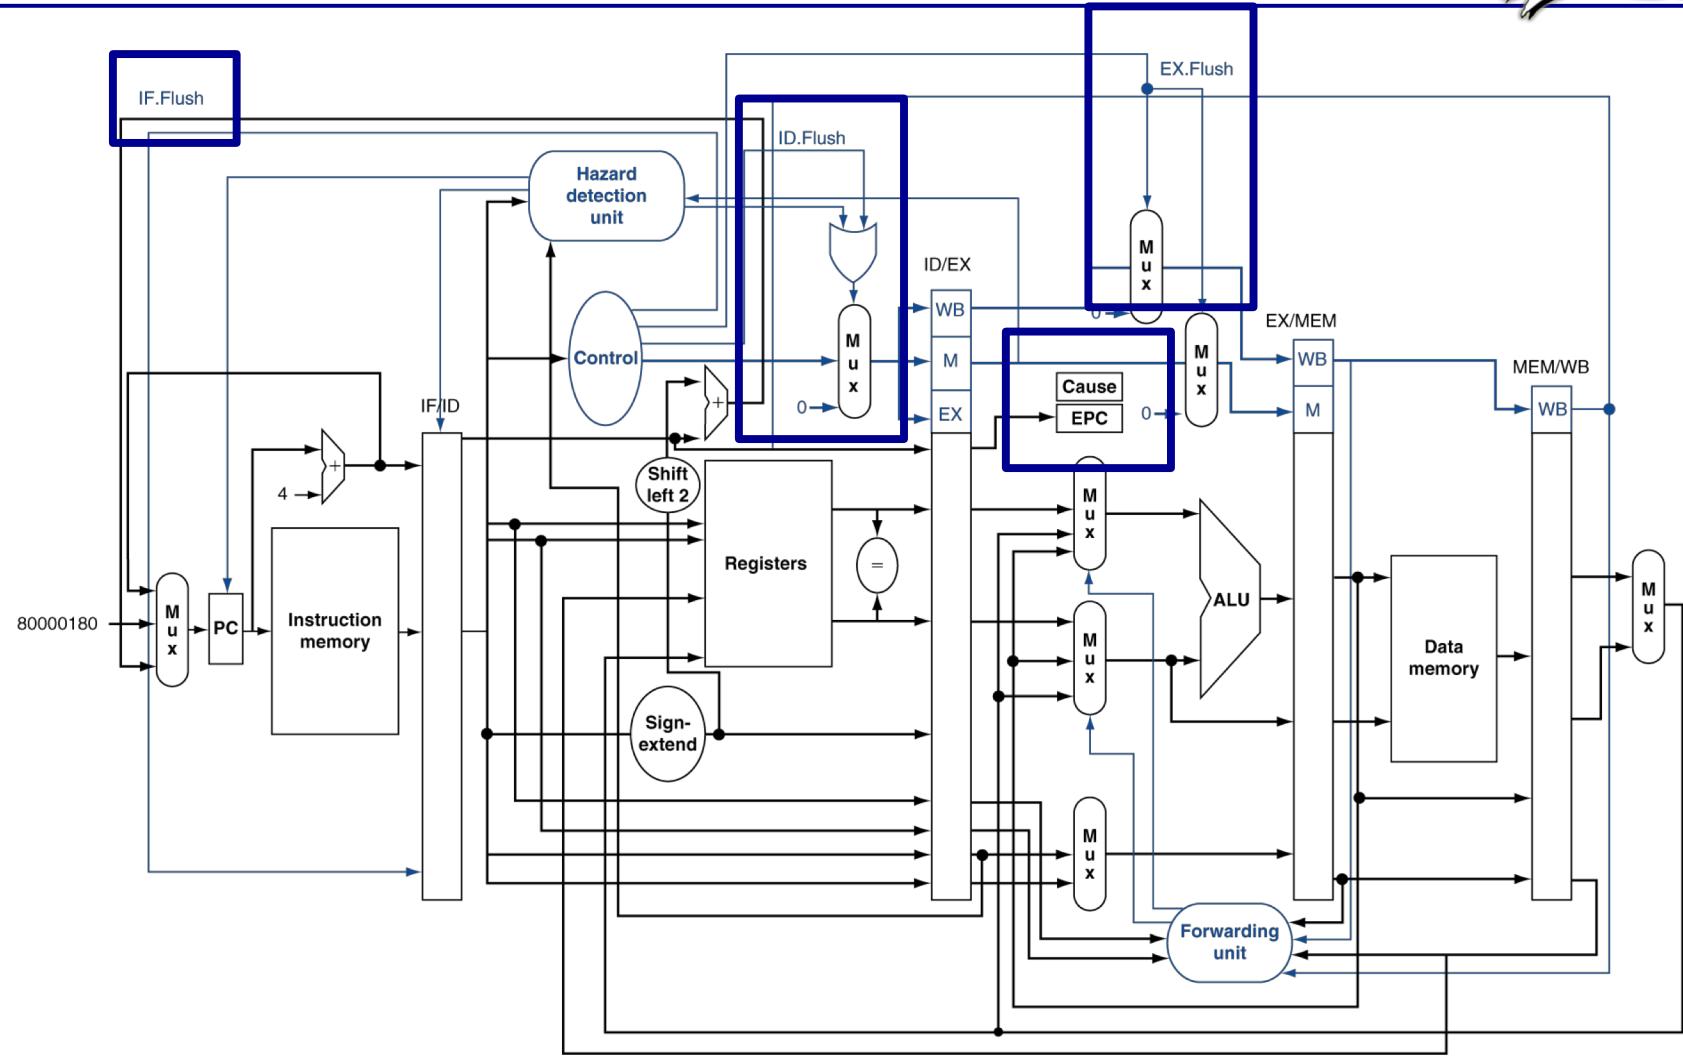
- Another form of control hazard
- Consider division by zero on add in EX stage
 - divi \$1, \$1, 0
 - Prevent \$1 from being overwritten with – what?
 - Complete previous instructions
 - Nullify subsequent instructions
 - Set Cause and EPC register values
 - Transfer control to handler
- Similar to mispredicted branch
 - Uses much of the same hardware

Nullifying Instructions



- Nullify = turn an instruction into a nop (or bubble)
 - Reset its RegWrite and MemWrite signals
 - Does not affect the state of the system
 - Resets its branch and jump signals
 - Does not cause unexpected flow control
 - Mark that it should not raise any exceptions of its own
 - Does not cause unexpected flow control
 - Let it flow down the pipeline

5-stage Pipeline with Exceptions



Exception Example



■ Exception on add in

```
40      sub    $11, $2, $4
44      and    $12, $2, $5
48      or     $13, $2, $6
4C      add    $1, $2, $1
50      s1t    $15, $6, $7
54      lw     $16, 50($7)
```

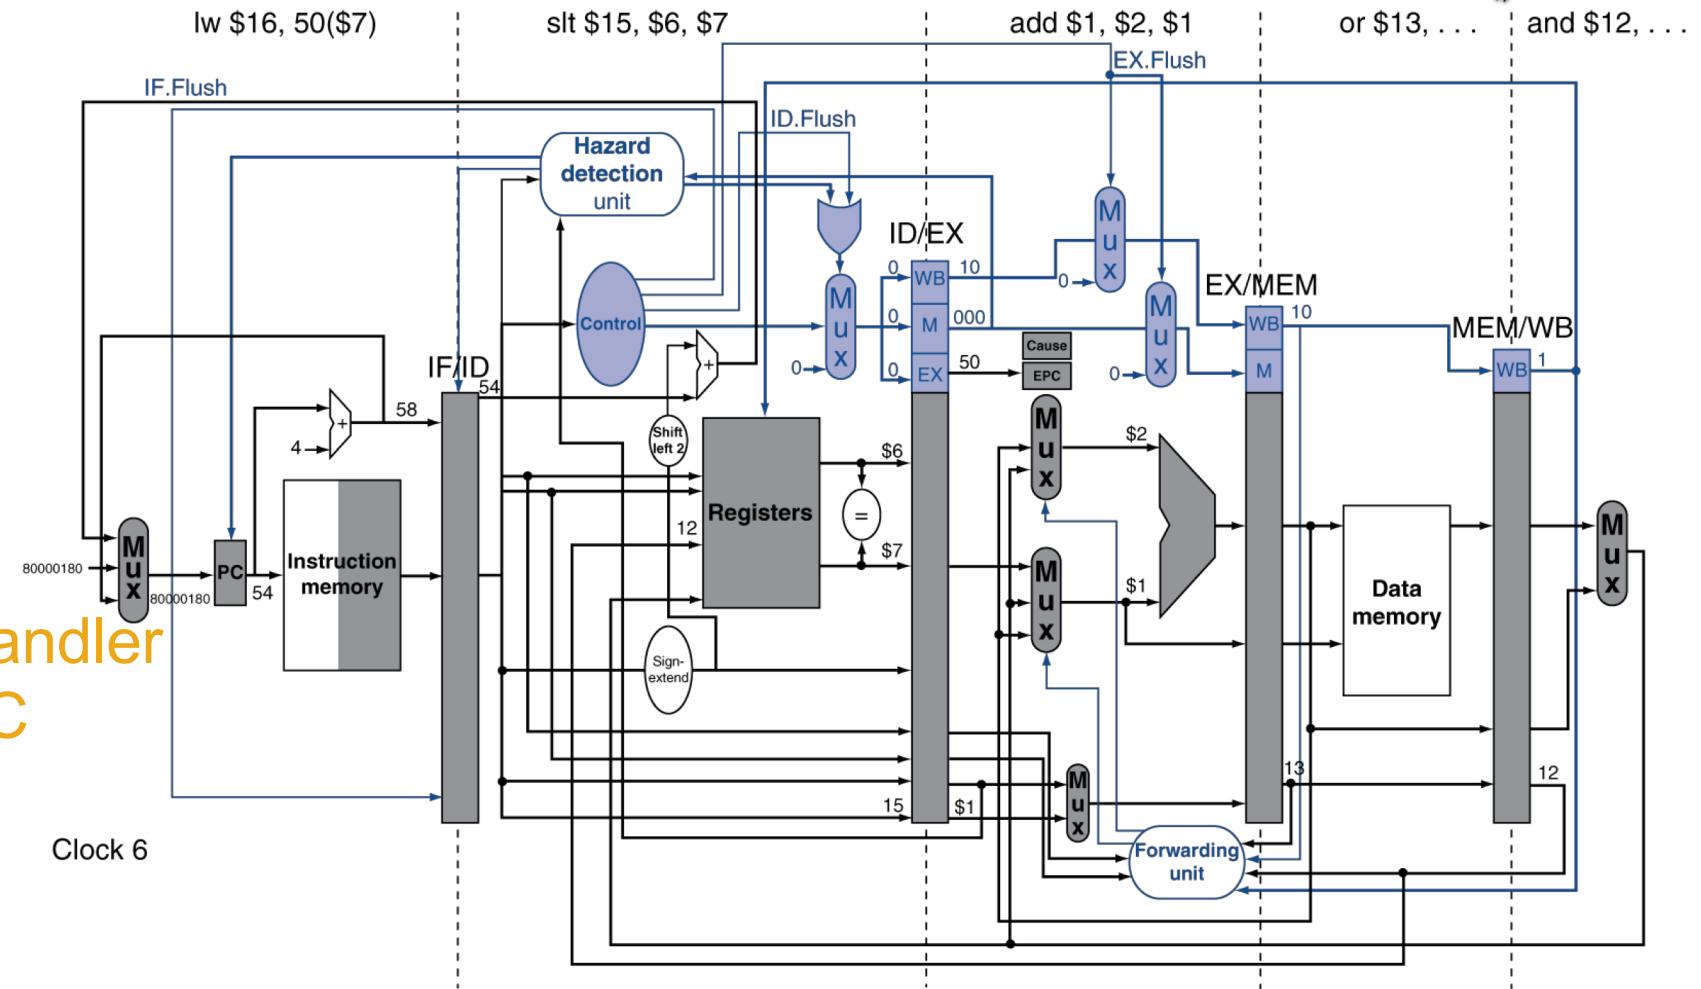
...

■ Handler

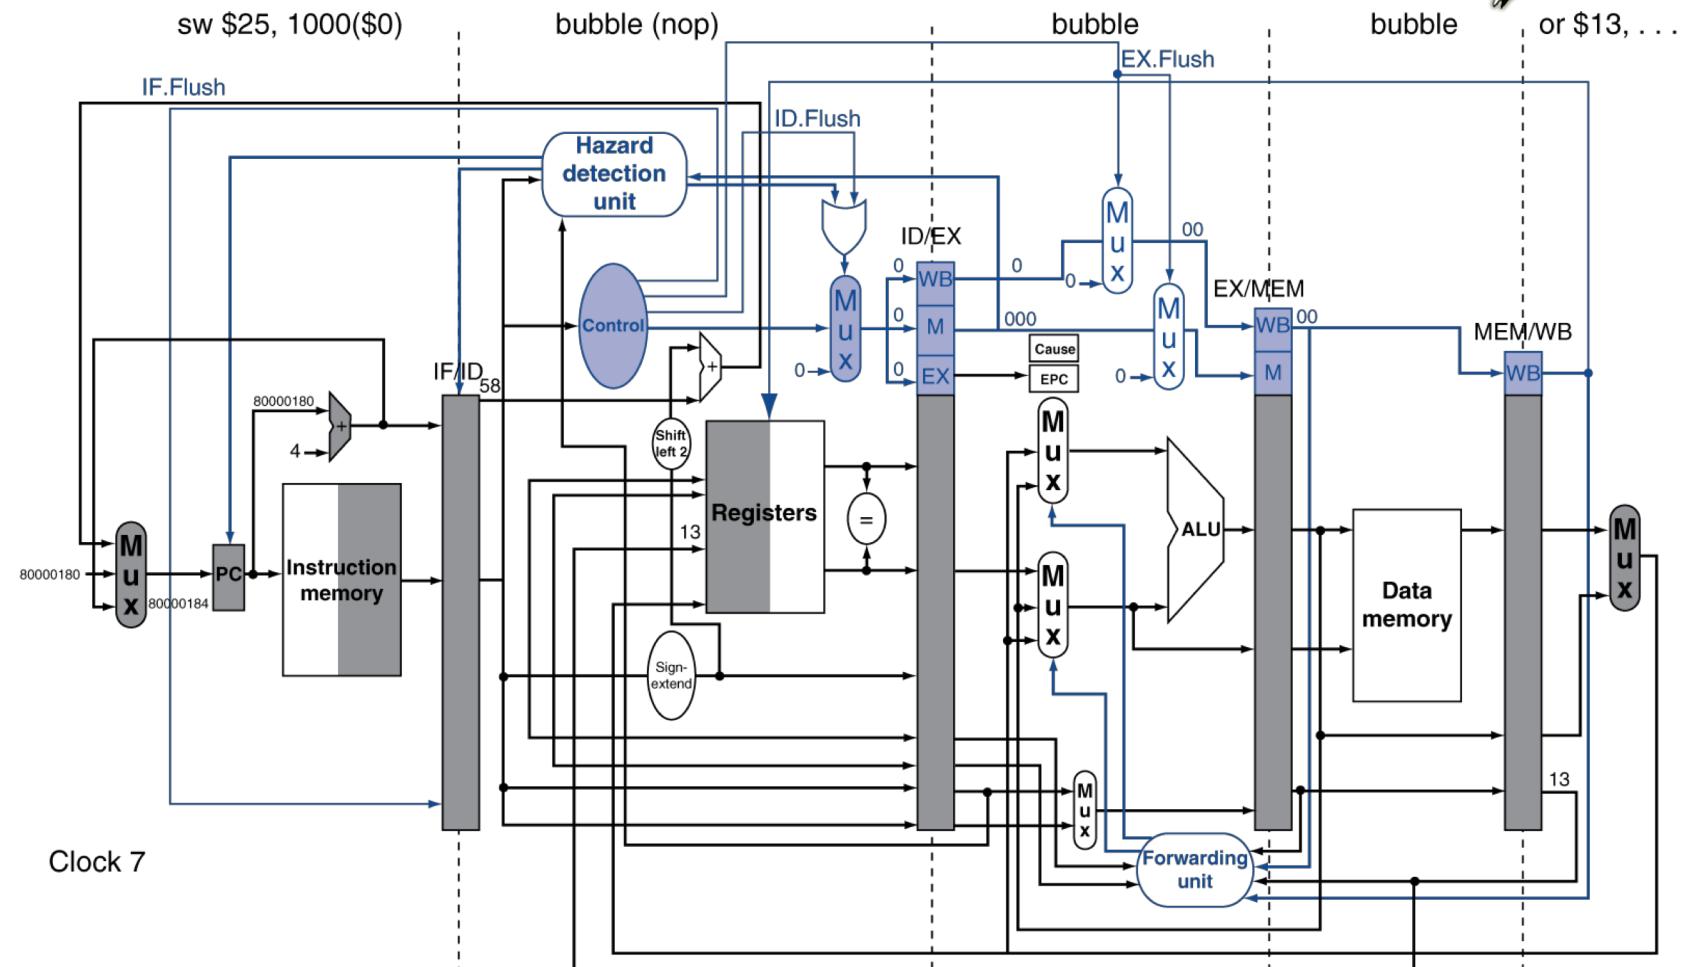
```
80000180      sw    $25, 1000($0)
80000184      sw    $26, 1004($0)
```

...

Exception Example



Exception Example



Dealing with Multiple Exceptions



- Pipelining overlaps multiple instructions
 - Could have multiple exceptions at once
- Simple approach: deal with exception in program order
 - Single commit point (MEM or WB stage)
 - Flush subsequent instructions
 - Necessary for “precise” exceptions

Beyond the 5-stage Pipeline



Desirable properties

Higher clock frequency

$CPI_{base} < 1$

Avoid in-order stalls

When instruction stalls, independent instructions behind it stall

Avoid stalls due to branches

How do we go about it?

Modern Processors



High clock frequency → deep pipelines

10 – 20 stages

$CPI_{base} < 1 \rightarrow$ superscalar pipelines

Launch 2, 3, or 4 instructions every cycle

Avoid in-order stalls → out-of-order execution

Re-order instructions based on dependencies

Avoid stalls due to branches → branch prediction

Keep history about direction and target of branches

Why Does it Work: Instruction-level Parallelism



Independence among instructions

Example with ILP

add \$t0, \$t1, \$t2 Most programs have ILP

or \$t3, \$t1, \$t2 But it is not infinite

sub \$t4, \$t1, \$t2 Scope also matters a lot

and \$t5, \$t1, \$t2

Example with no ILP

add \$t0, \$t1, \$t2

or \$t3, \$t0, \$t2

sub \$t4, \$t3, \$t2

and \$t5, \$t4, \$t2

Question



How much ILP is there in common programs?

Multiple Instruction Issue (Superscalar)



Fetch and execute multiple instructions per cycle => CPI < 1

Example: fetch 2 instructions/clock cycle

Dynamically decide if they can go down the pipe together or not (hazard)

Pipe Stages

	IF	ID	EX	MEM	WB		
	IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB	
Ideal CPI = 0.5		IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB
			IF	ID	EX	MEM	WB

Resources: double amount of hardware (FUs, Register file ports)

Issues: hazards, branch delay, load delay

Modern processors: Intel x86: 4-way, ARM: 6-way/8-way

Deeper Pipelining



Increase number of pipeline stages

Fewer levels of logic per pipe stage

Higher clock frequency

Example 9-stage pipeline

Pipe Stages

IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB
	IF1	IF2	ID	EX	MEM1	MEM2	MEM3 WB

Issues: branch delay, load delay: $CPI = 1.4 - 2.0$, cost of pipeline registers

Modern pipelines

Pentium: 5 stages

Original Pentium Pro: 12 stages

Pentium 4: 31 stages, 3+ GHz

Tejas: 40-50 stages, 10 GHz, cancelled!

Core 2: 14 stages

Dynamic Scheduling



Execute instructions out-of-order

Fetch multiple instructions per cycle using branch prediction

Figure out which are independent and execute them in parallel

Example

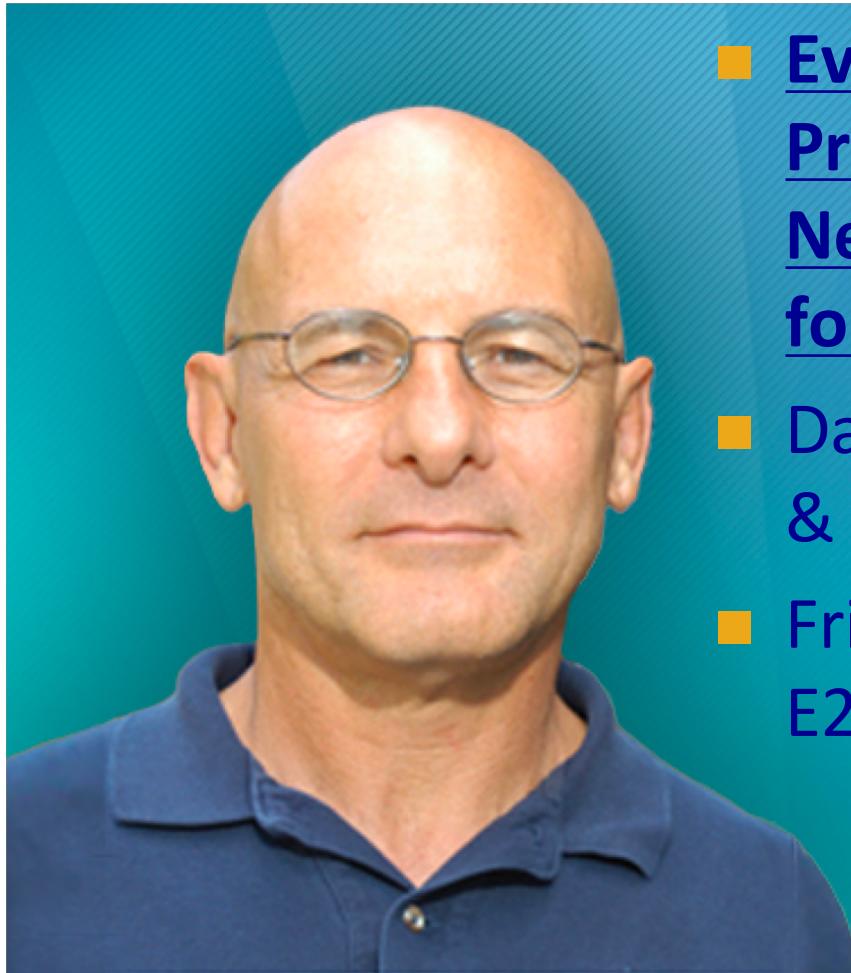
```
add    $t0, $t1, $t2  
or     $t3, $t0, $t2  
sub    $t0, $t1, $t2  
and    $t5, $t0, $t2
```

Superscalar + Dynamic scheduling

add \$t0, \$t1, \$t2	sub \$t0, \$t1, \$t2
or \$t3, \$t0, \$t2	and \$t5, \$t0, \$t2

What's wrong with this?

Announcements

A portrait photograph of David Patterson, a middle-aged man with a shaved head and glasses, wearing a blue polo shirt. He is positioned on the left side of the slide, with a teal gradient background behind him.

- Evaluation of the Tensor Processing Unit: A Deep Neural Network Accelerator for the Datacenter
- David Patterson, UC Berkeley & Google
- Friday, February 8, 2:40 PM, E2-180