

CMPE110 Lecture 15

Exceptions

Heiner Litz

<https://canvas.ucsc.edu/courses/19290>

Announcements



- No Quiz on Friday
- Midterm: Wed 02/13
 - 1 double sided page of notes
 - RISC-V Green card/cheat sheet part of the exam

Review



Modern Processors



High clock frequency → deep pipelines

10 – 20 stages

$CPI_{base} < 1 \rightarrow$ superscalar pipelines

Launch 2, 3, or 4 instructions every cycle

Avoid in-order stalls → out-of-order execution

Re-order instructions based on dependencies

Avoid stalls due to branches → branch prediction

Keep history about direction and target of branches

Instruction-Level Parallelism

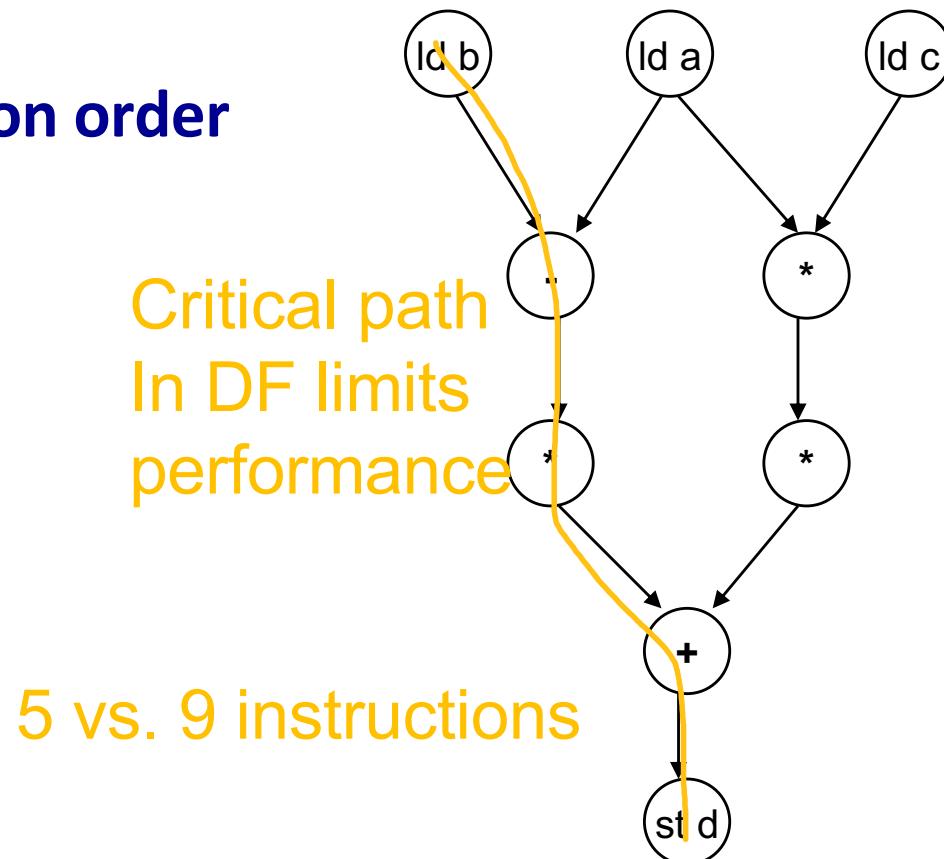


$$D = 3(a - b) + 7ac$$

■ Data-flow execution order

■ Sequential execution order

ld a
ld b
sub a-b
mul 3(a-b)
ld c
mul ac
mul 7ac
add 3(a-b)+7ac
st d



Question



How much ILP is there in common programs?

Limits on Instruction Level Parallelism (ILP)



Weiss and Smith [1984]	1.58
Sohi and Vajapeyam [1987]	1.81
Tjaden and Flynn [1970]	1.86
Tjaden and Flynn [1973]	1.96
Uht [1986]	2.00
Smith et al. [1989]	2.00
Jouppi and Wall [1988]	2.40
Johnson [1991]	2.50
Acosta et al. [1986]	2.79
Wedig [1982]	3.00
Butler et al. [1991]	5.8
Melvin and Patt [1991]	6
Wall [1991]	7
Kuck et al. [1972]	8
Riseman and Foster [1972]	51
Nicolau and Fisher [1984]	90

- These studies made very different assumptions!
- Any guesses on the differences?

Multiple Instruction Issue (Superscalar)



Fetch and execute multiple instructions per cycle => CPI < 1

Example: fetch 2 instructions/clock cycle

Dynamically decide if they can go down the pipe together or not (hazard)

Pipe Stages

	IF	ID	EX	MEM	WB		
	IF	ID	EX	MEM	WB		
		IF	ID	EX	MEM	WB	
Ideal CPI = 0.5		IF	ID	EX	MEM	WB	
			IF	ID	EX	MEM	WB
			IF	ID	EX	MEM	WB

Resources: double amount of hardware (FUs, Register file ports)

Issues: hazards, branch delay, load delay

Modern processors: Intel x86: 4-way, ARM: 6-way/8-way



Deeper Pipelining

Increase number of pipeline stages

Fewer levels of logic per pipe stage

Higher clock frequency

Example 9-stage pipeline

Pipe Stages

IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB	
	IF1	IF2	ID	EX	MEM1	MEM2	MEM3	WB

Issues: branch delay, load delay: $CPI = 1.4 - 2.0$, cost of pipeline registers

Modern pipelines

Pentium: 5 stages

Original Pentium Pro: 12 stages

Pentium 4: 31 stages, 3+ GHz

Tejas: 40-50 stages, 10 GHz, cancelled!

Core 2: 14 stages



Dynamic Scheduling

Execute instructions out-of-order

Fetch multiple instructions per cycle using branch prediction

Figure out which are independent and execute them in parallel

Example

add	\$t0	,	\$t1,	\$t2
or	\$t3,	\$t0,	\$t2	
sub	\$t0,	\$t1,	\$t2	
and	\$t5,	\$t0,	\$t2	

Superscalar + Dynamic scheduling

add	\$t0,	\$t1,	\$t2
or	\$t3,	\$t0,	\$t2

sub	\$t0,	\$t1,	\$t2
and	\$t5,	\$t0,	\$t2

What's wrong with this?

With OOO WAW and RAW hazards matter!



Register Renaming

Rename (map) architectural registers to physical registers in decode stage to get rid of false dependencies

add \$t0, \$t1, \$t2

or \$t3, \$t0, \$t2

sub \$t0, \$t1, \$t2

and \$t5, \$t0, \$t2

Superscalar + Dynamic scheduling + register renaming

add \$t0_A, \$t1, \$t2 sub \$t0_B, \$t1, \$t2

or \$t3, \$t0_A, \$t2 and \$t5, \$t0_B, \$t2

Need more physical registers than architectural registers

Physical registers are automatically recycled

How do we know which instructions can be scheduled?



- Need to represent the data flow graph in hardware

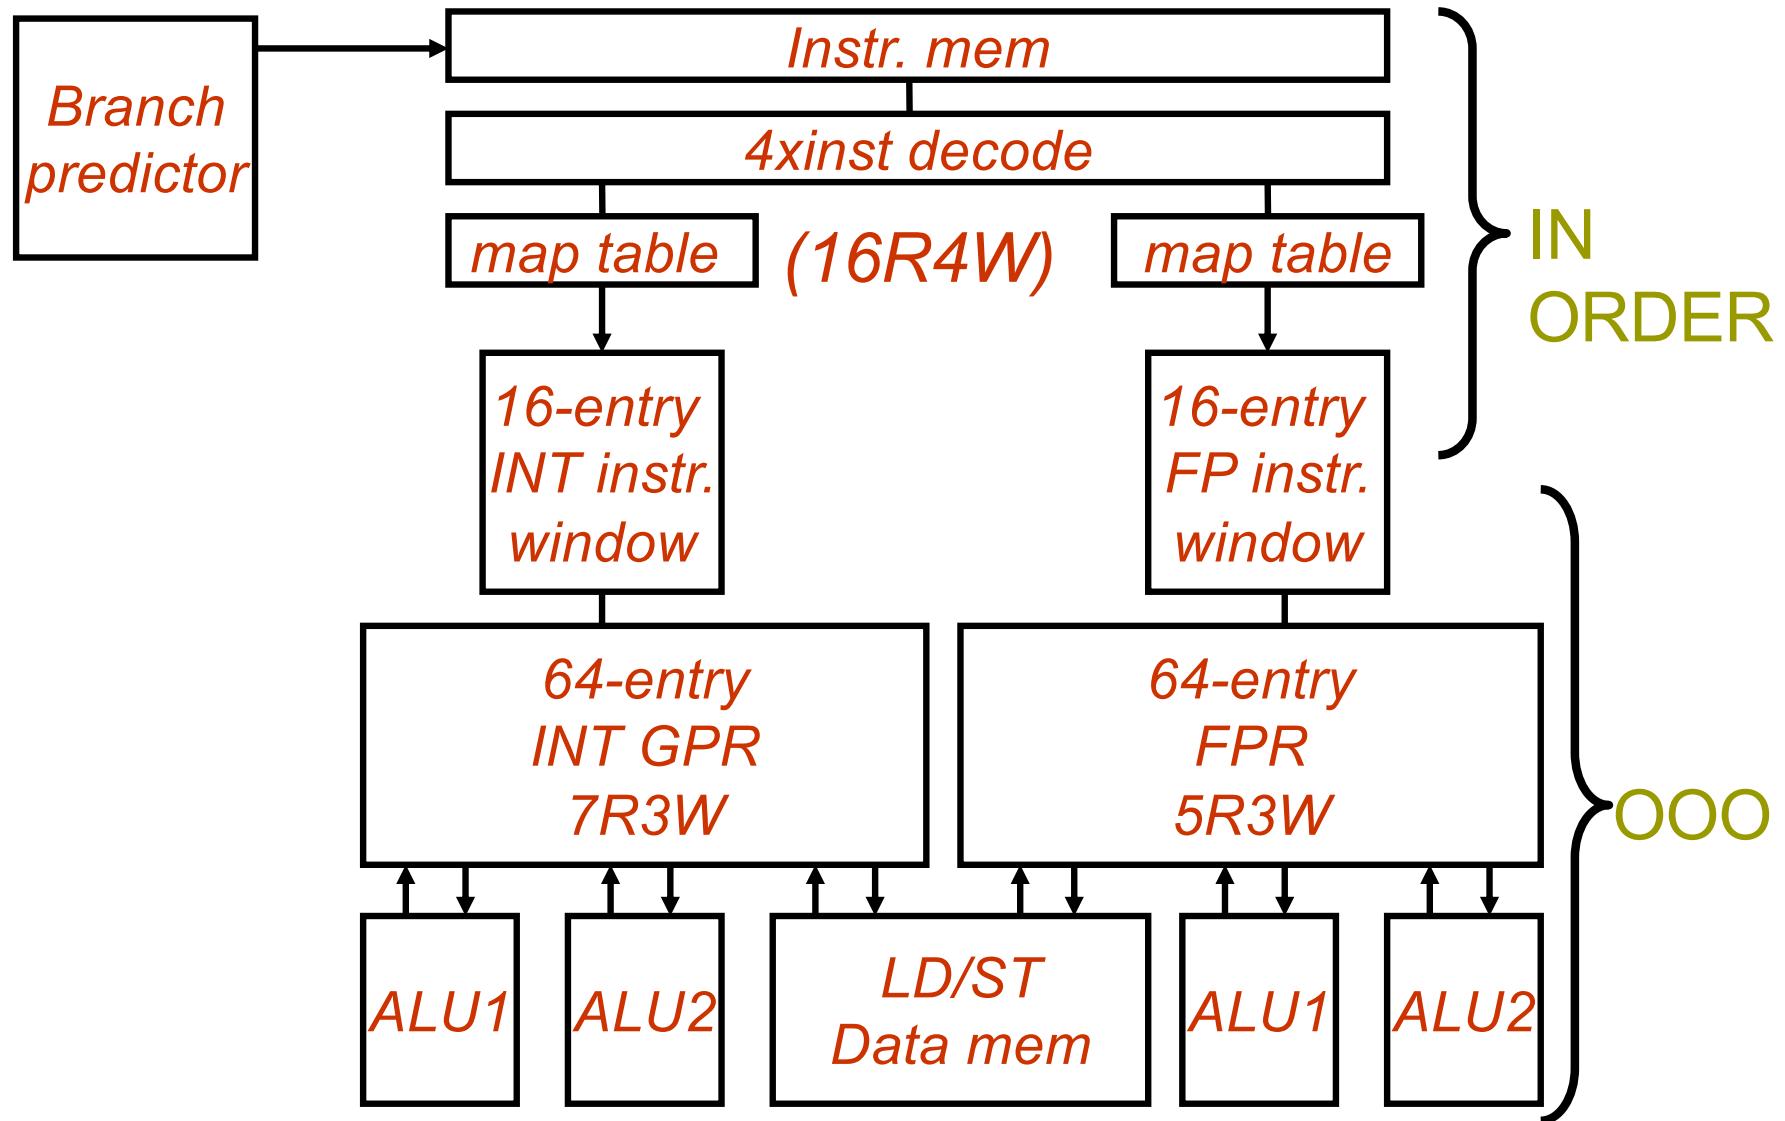


Simple Scoreboarding

- Scoreboard: a bit-array, 1-bit for each GPR
 - if the bit is not set, the register has valid data
 - if the bit is set, the register has stale data
 - i.e. some outstanding instruction is going to change it
- Dispatch in order: $RD \leftarrow Fn(RS, RT)$
 - if $SB[RS]$ or $SB[RT]$ is set \Rightarrow RAW, stall
 - if $SB[RD]$ is set is set \Rightarrow WAW, stall
 - else dispatch to FU, set $SB[RD]$
 - Assuming a functional unit is available (structural hazard)
- Complete out-of-order
 - update GPR[RD], clear $SB[RD]$

Can be avoided
with renaming

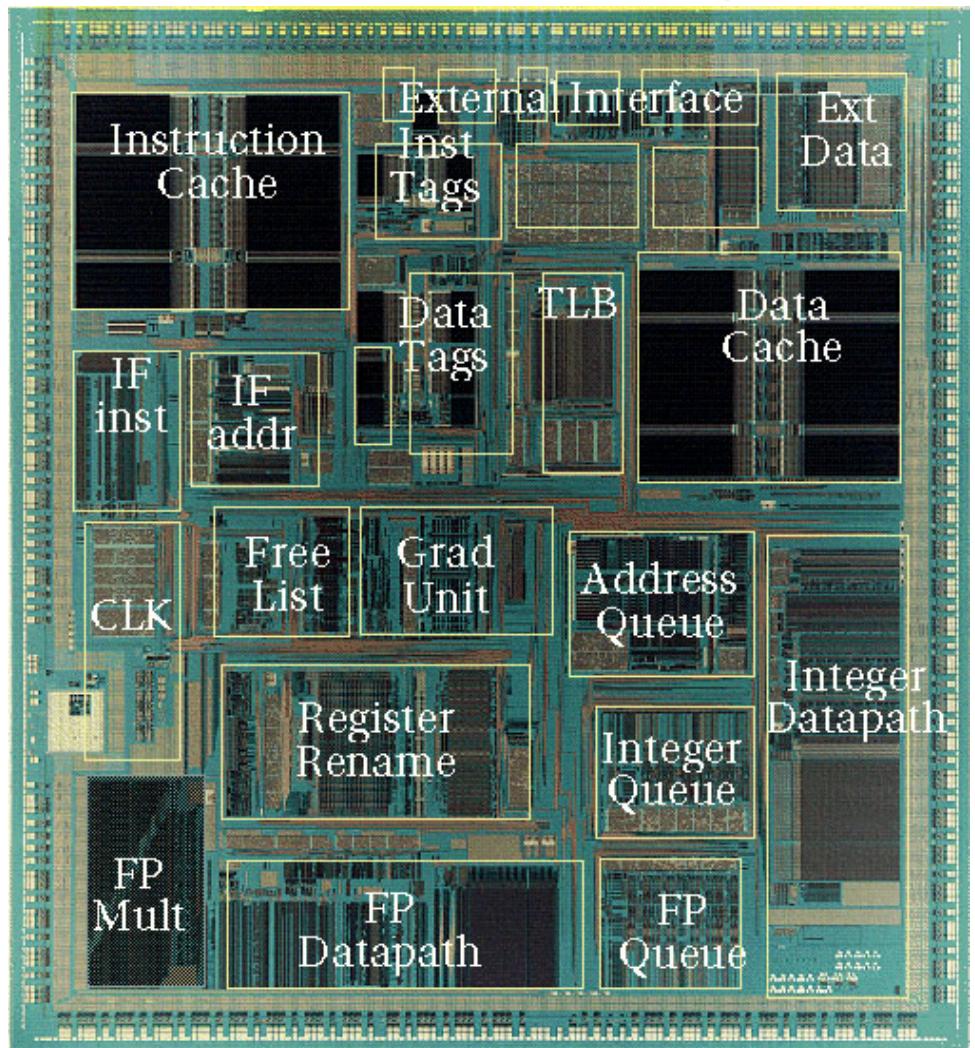
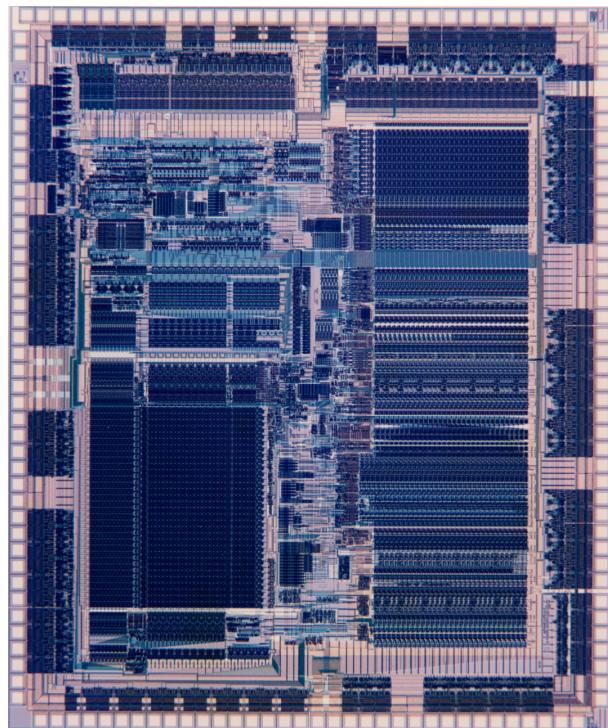
Dynamic Scheduling in a Modern 000 MIPS R10000





R2000 vs. R10000

- R2000: mid-80's, 32b 5-stage integer pipeline only, 120 K transistors
- R10k: mid-90's, 64b out-of-order superscalar w/ integer & FPU, 64 KB cache, 6.8 M transistors





Limits of Advanced Pipelining

Limited ILP in real programs

Pipeline overhead

Branch and load delays exacerbated

Clock cycle timing limits

Limited branch prediction accuracy (85%-98%)

Even a few percent really hurts with long/wide pipes!

Memory inefficiency

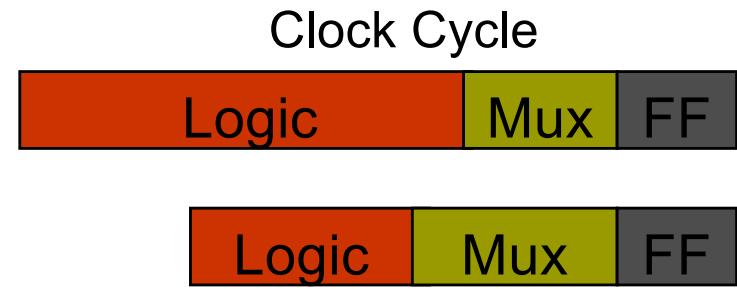
Load delays + # of loads/cycle

Caches (next lecture) are not perfect

Complexity of implementation

Expensive to design

Why does a superscalar, out-of-order processor burn more power and energy?



Caveats of Out-of-Order Processors



What can go wrong if I reorder instructions?

Tips

- Think of dependencies
- Think of forwarding
- Think of exceptions

Solutions?

Performance/Power Tradeoffs

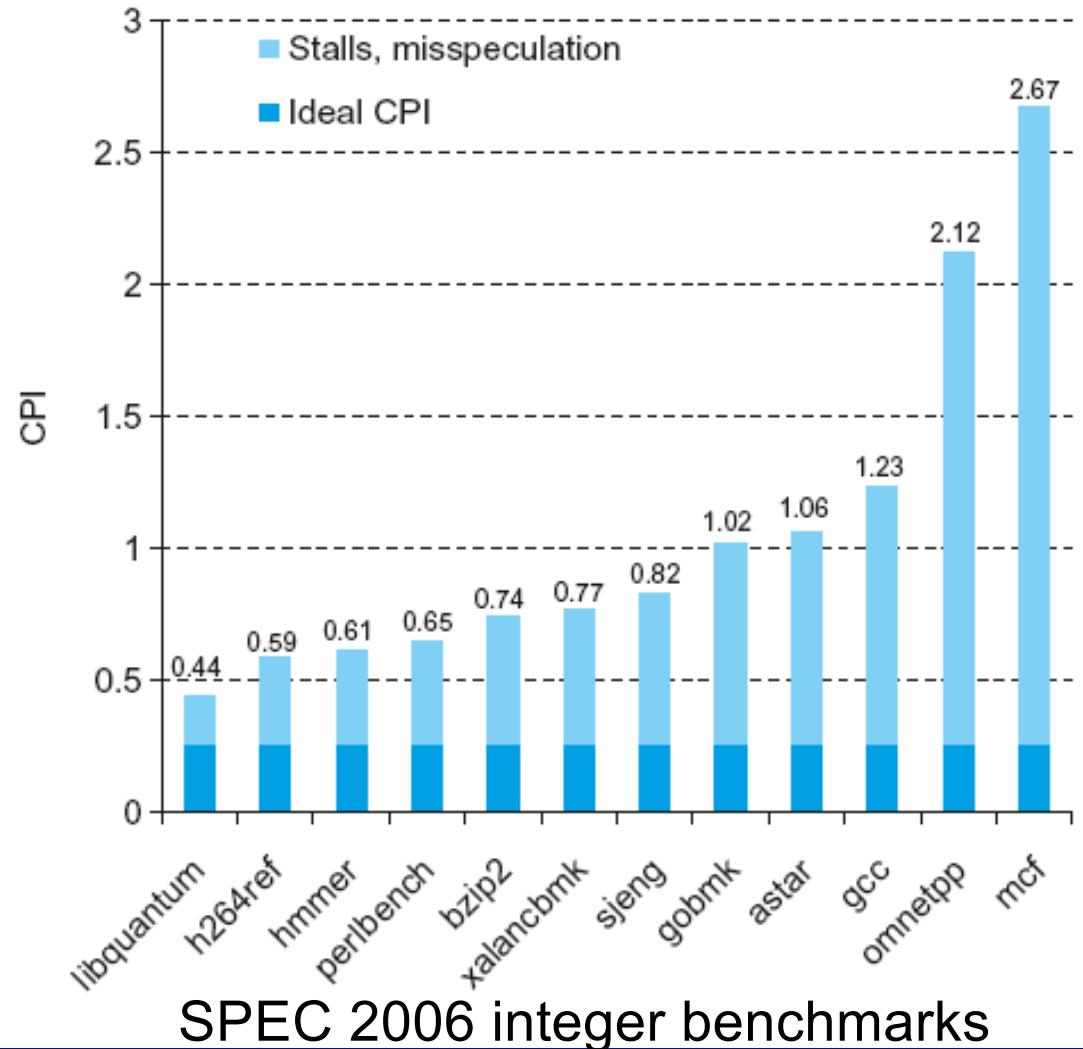


Processor	ARM A8	Intel Core i7 920
Market	Personal Mobile Device	Server, cloud
Thermal design power	2 Watts	130 Watts
Clock rate	1 GHz	2.66 GHz
Cores/Chip	1	4
Floating point?	No	Yes
Multiple issue?	Yes	Yes
Peak instructions/clock cycle	2	4
Pipeline stages	14	14
Pipeline schedule	Static in-order	Dynamic out-of-order with speculation
Branch prediction	2-level, simple	2-level, aggressive
1 st level caches/core	32 KiB I, 32 KiB D	32 KiB I, 32 KiB D
2 nd level caches/core	128-1024 KiB	256 KiB
3 rd level caches (shared)	-	2- 8 MB

Performance/Power Tradeoffs

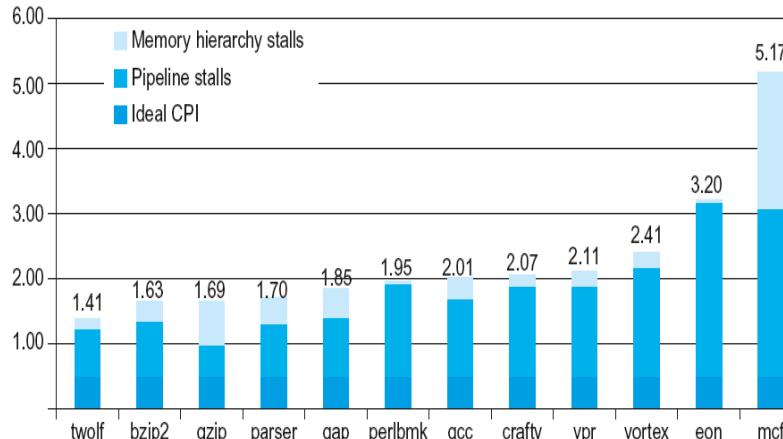


Intel i7 920



Minnespec Benchmarks

ARM A8



Impact of Branch Missprediction

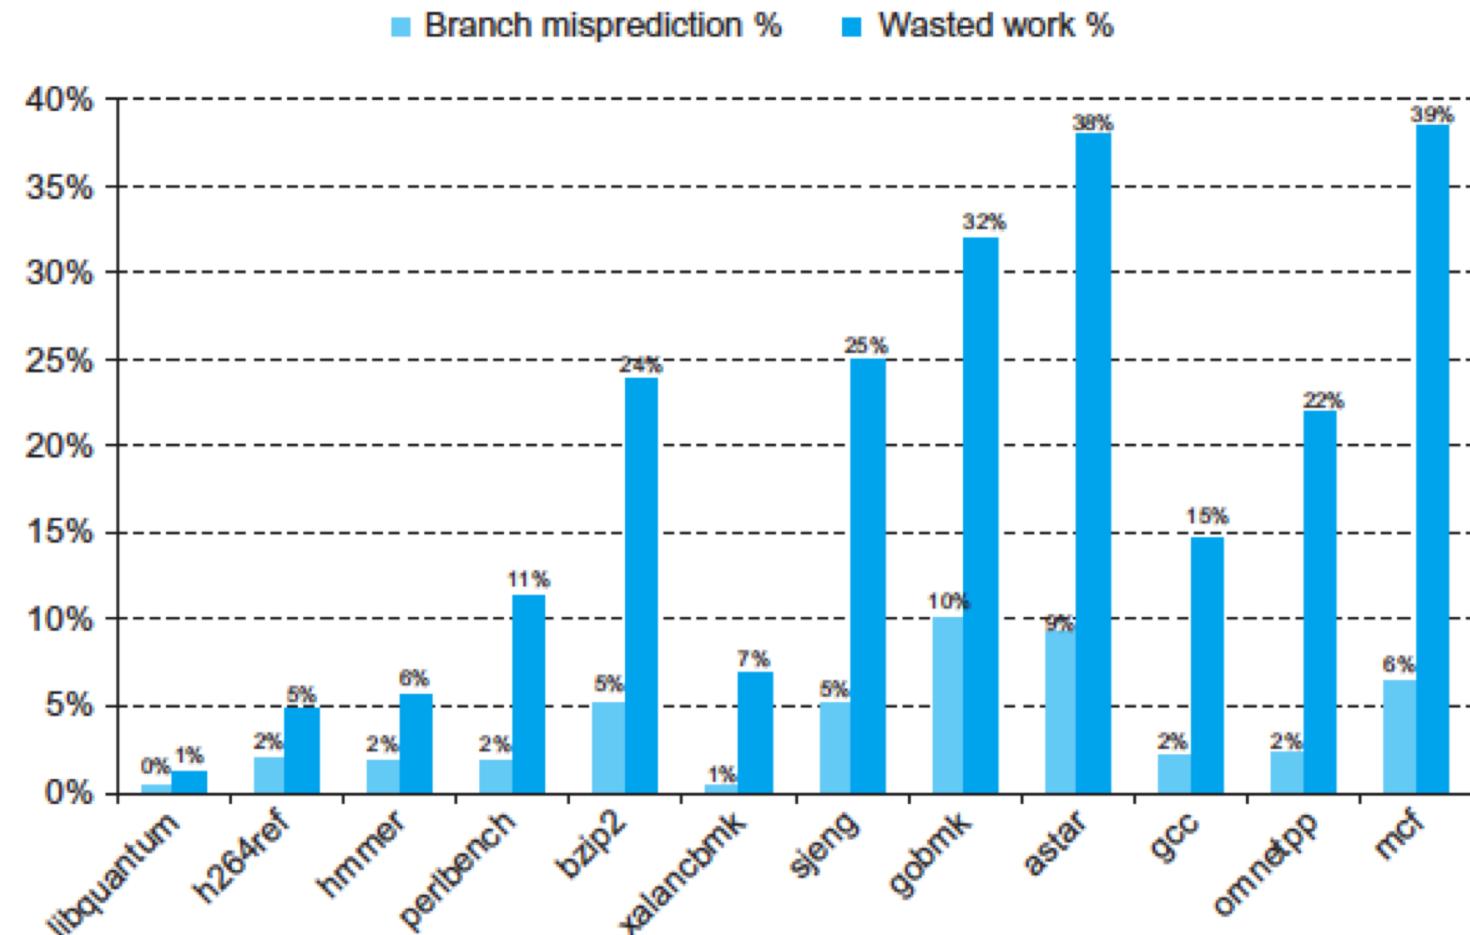


FIGURE 4.79 Percentage of branch mispredictions and wasted work due to unfruitful speculation of Intel Core i7 920 running SPEC2006 integer benchmarks.