

CMPE110 Lecture 18

Cache Coherency

Heiner Litz

<https://canvas.ucsc.edu/courses/19290>

Announcements



- Readings: 5.1-5.4, 5.9-5.10

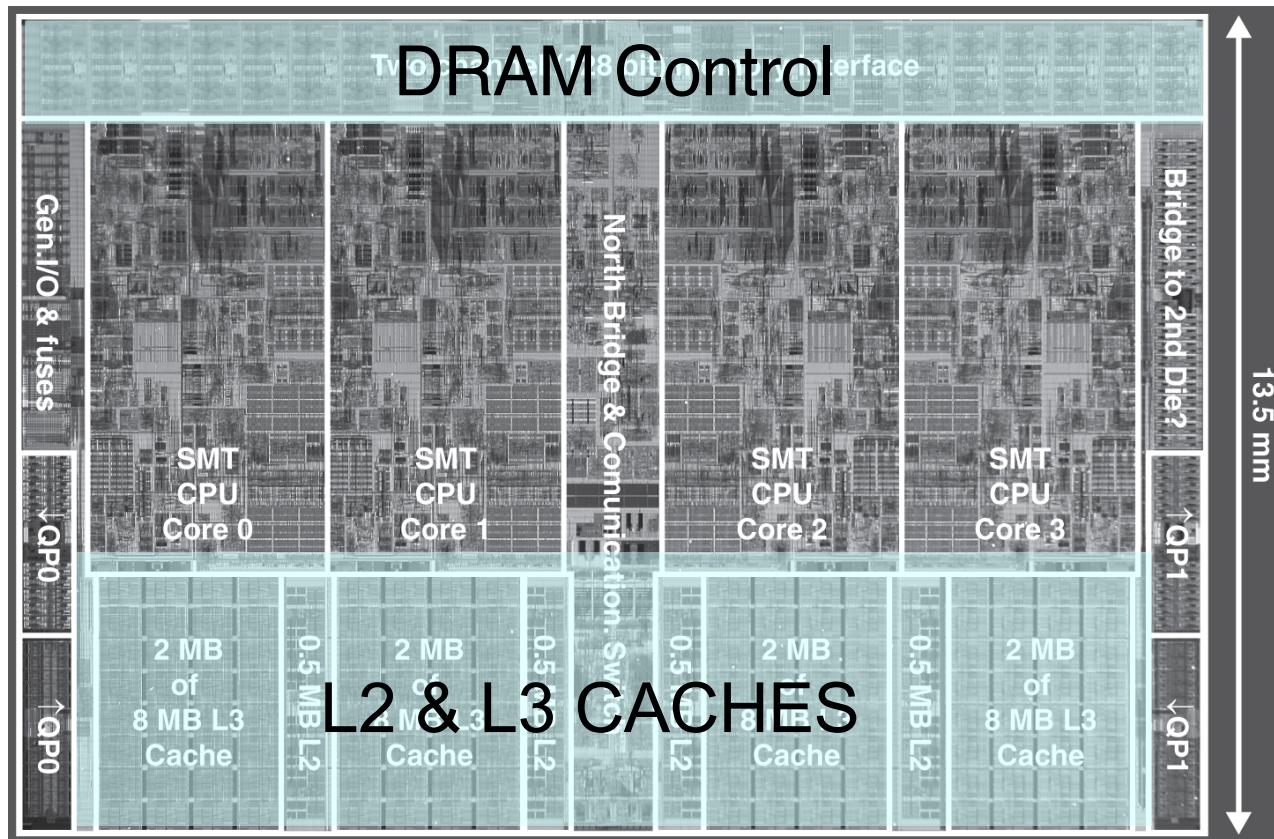
Review



Multilevel On-Chip Caches



Intel Nehalem 4-core processor



Per core:

- 32KB, 4-way L1 \$I
- 32KB, 8-way L1 \$D
- 256KB, 8-way L2

Shared

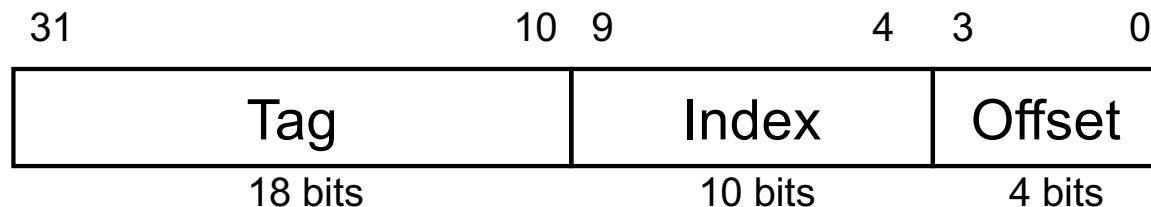
- 8 MB, 16-way L3



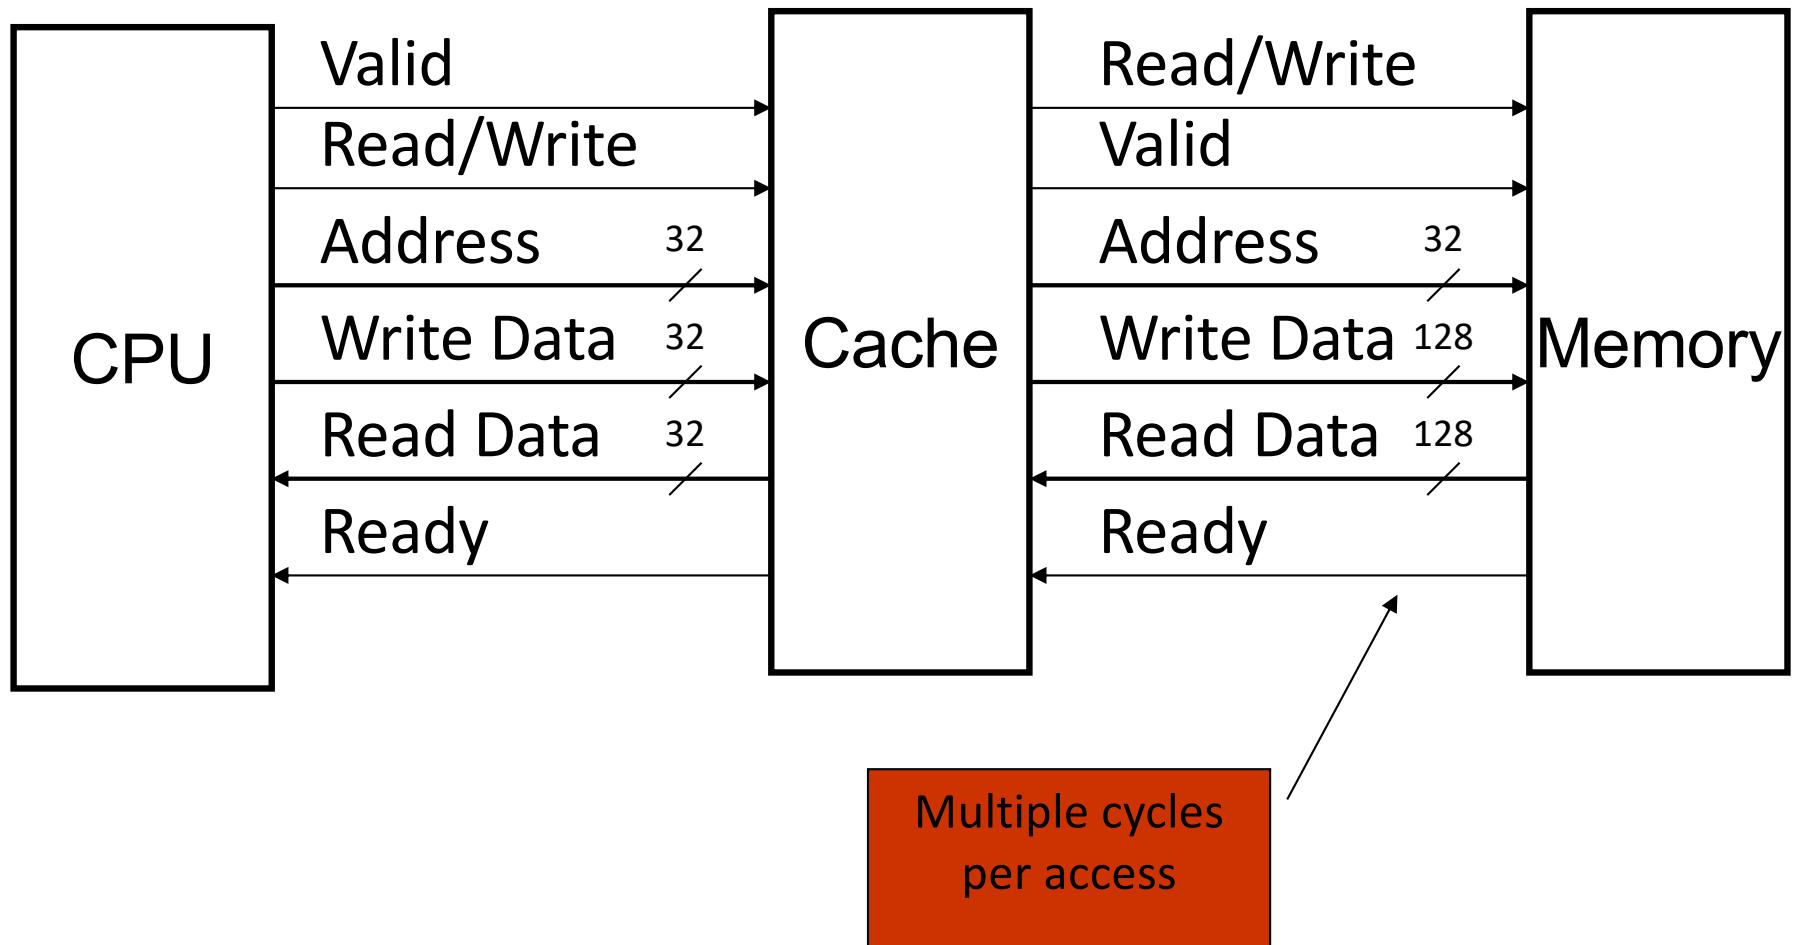
Cache Implementation: Control

Example cache characteristics

- Direct-mapped, write-back, write allocate
- Block size: 4 words (16 bytes)
- Cache size: 16 KB (1024 blocks)
- 32-bit byte addresses
- Valid bit and dirty bit per block
- Blocking cache
 - CPU waits until access is complete



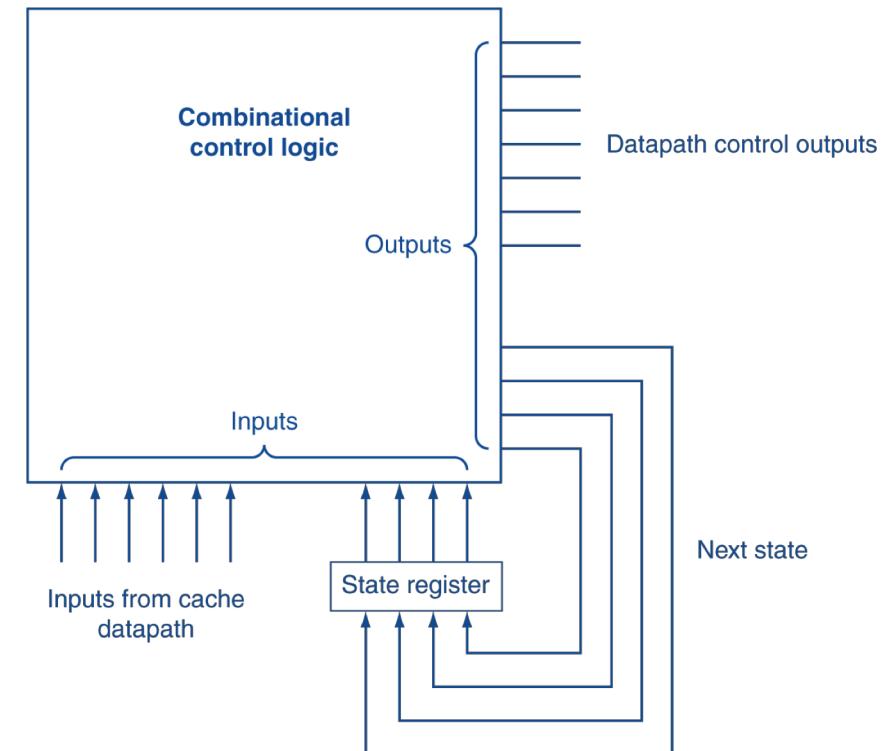
Interface Signals



Reminder: Finite State Machines

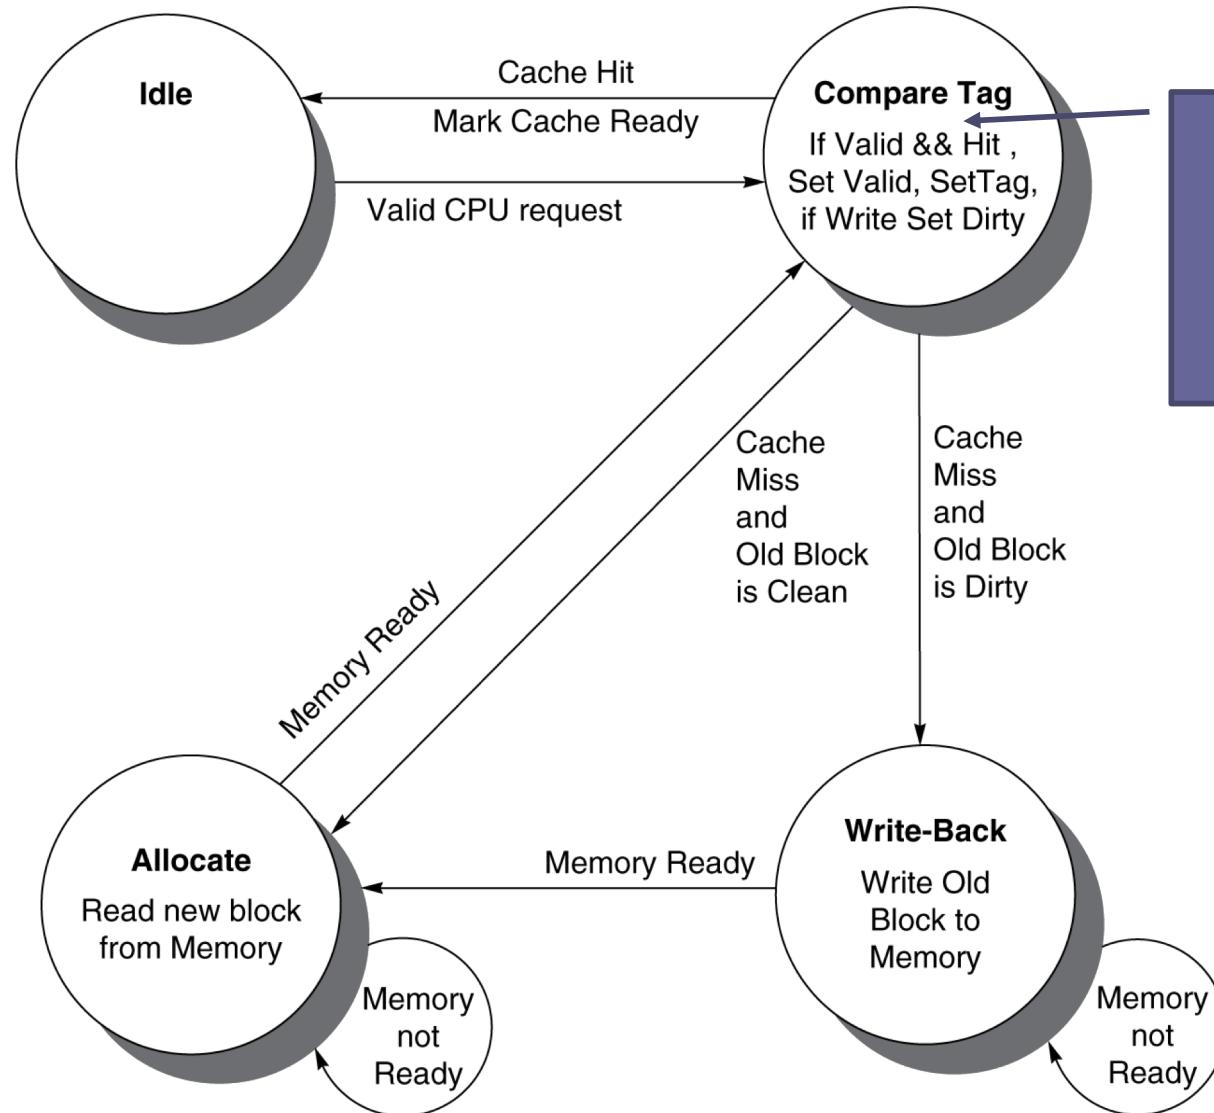


- Use an FSM to sequence control steps
- Set of states, transition on each clock edge
 - State values are binary encoded
 - Current state stored in a register
 - Next state
 $= fn(\text{current state}, \text{current inputs})$



- Control output signals $= fn(\text{current state})$

Cache Controller FSM - WRITE



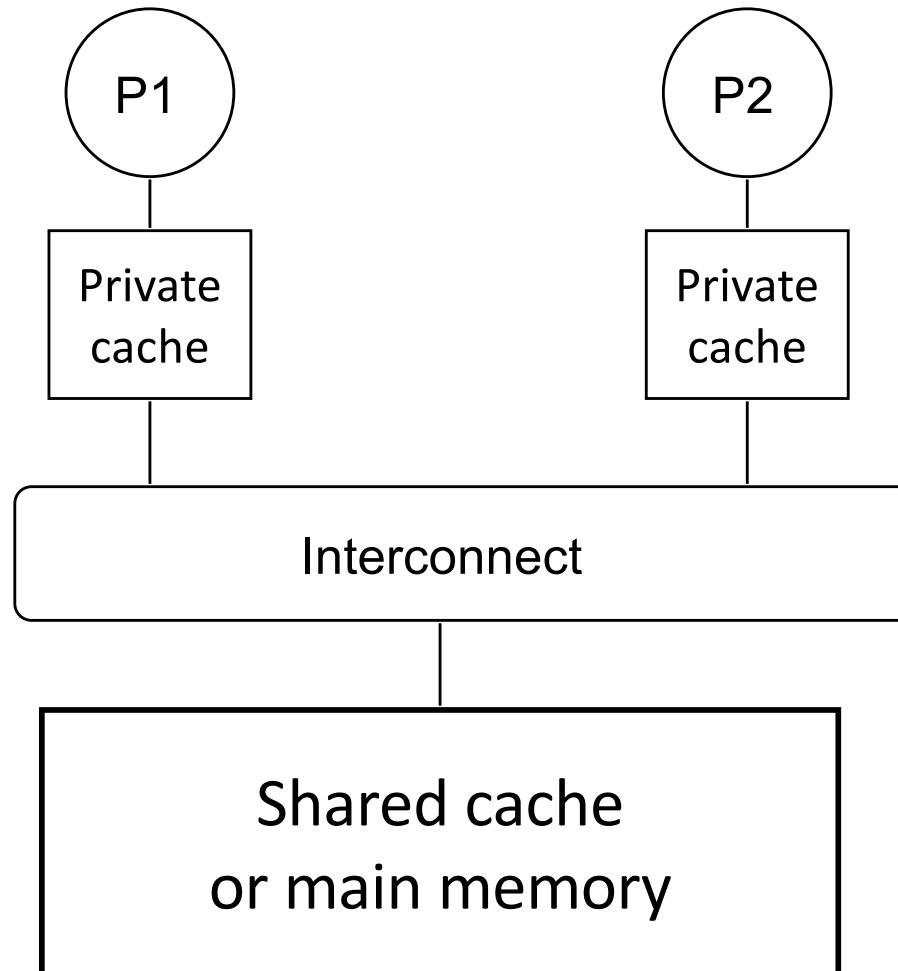
Could partition
into separate
states to reduce
clock cycle time

Multi-core & Caches



- See any issues?

Cache Coherence Problem





Cache Coherence

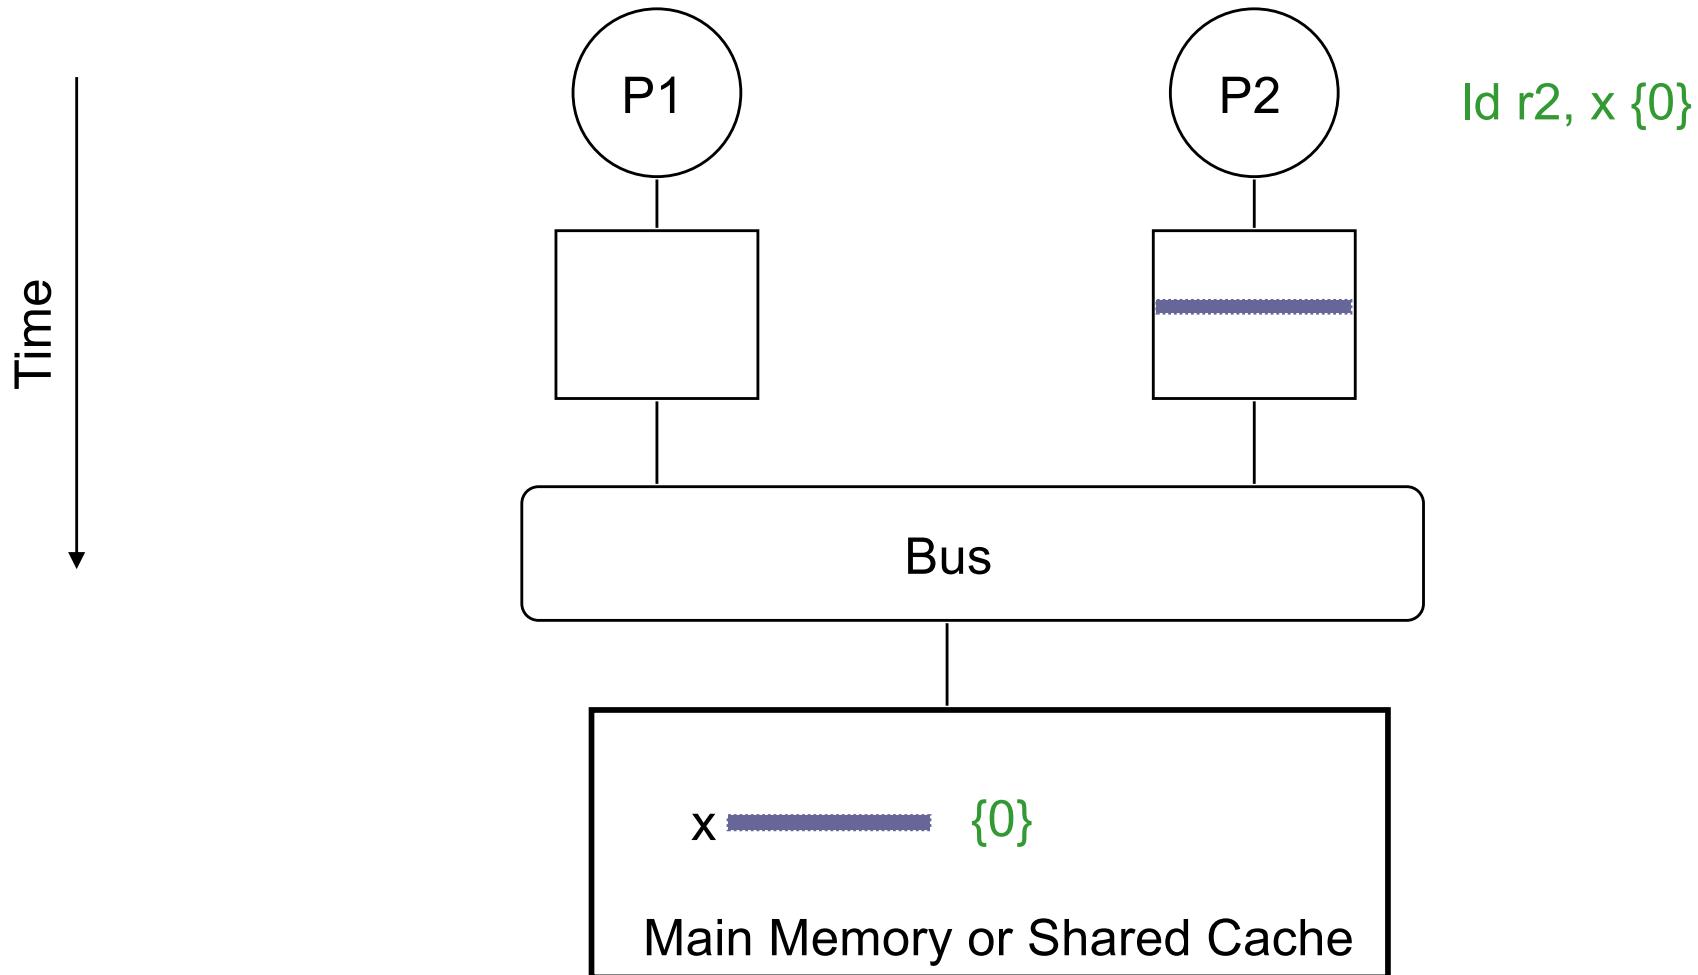
■ Informally

- Reads return most recently written value

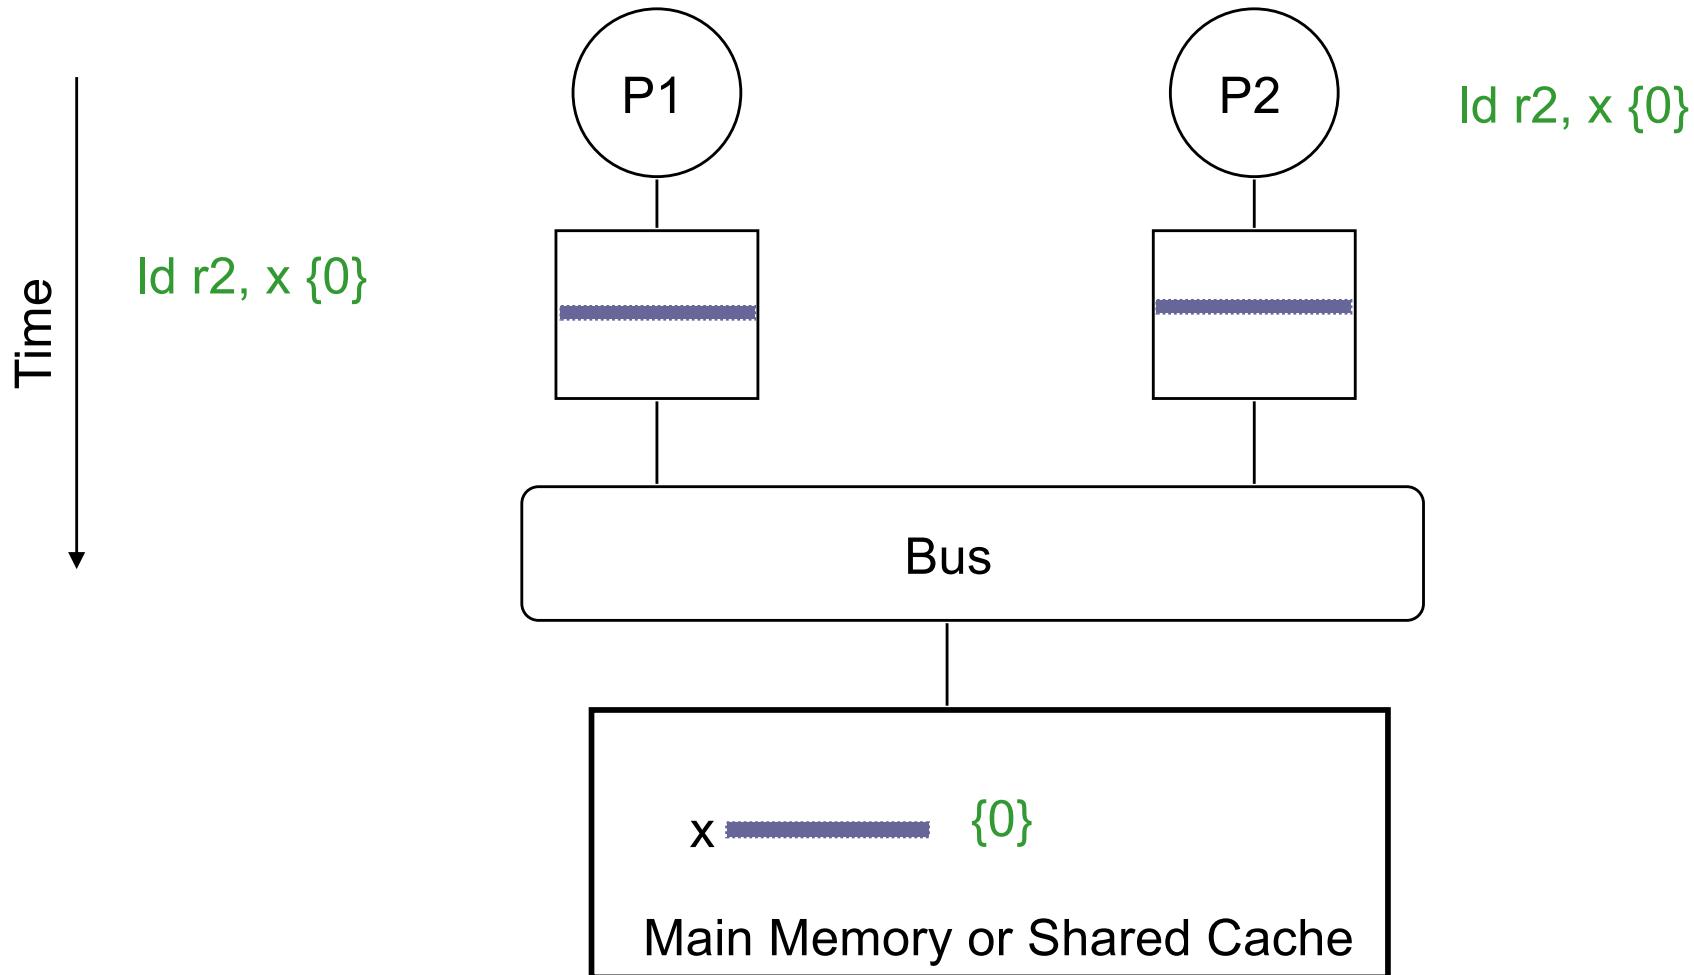
■ Formally

1. P writes X; P reads X (no intervening writes)
⇒ read returns written value
2. P1 writes X; P2 reads X (sufficiently later)
⇒ read returns written value
3. P1 writes X, P2 writes X
⇒ all processors see writes in the same order
⇒ End up with the same final value for X

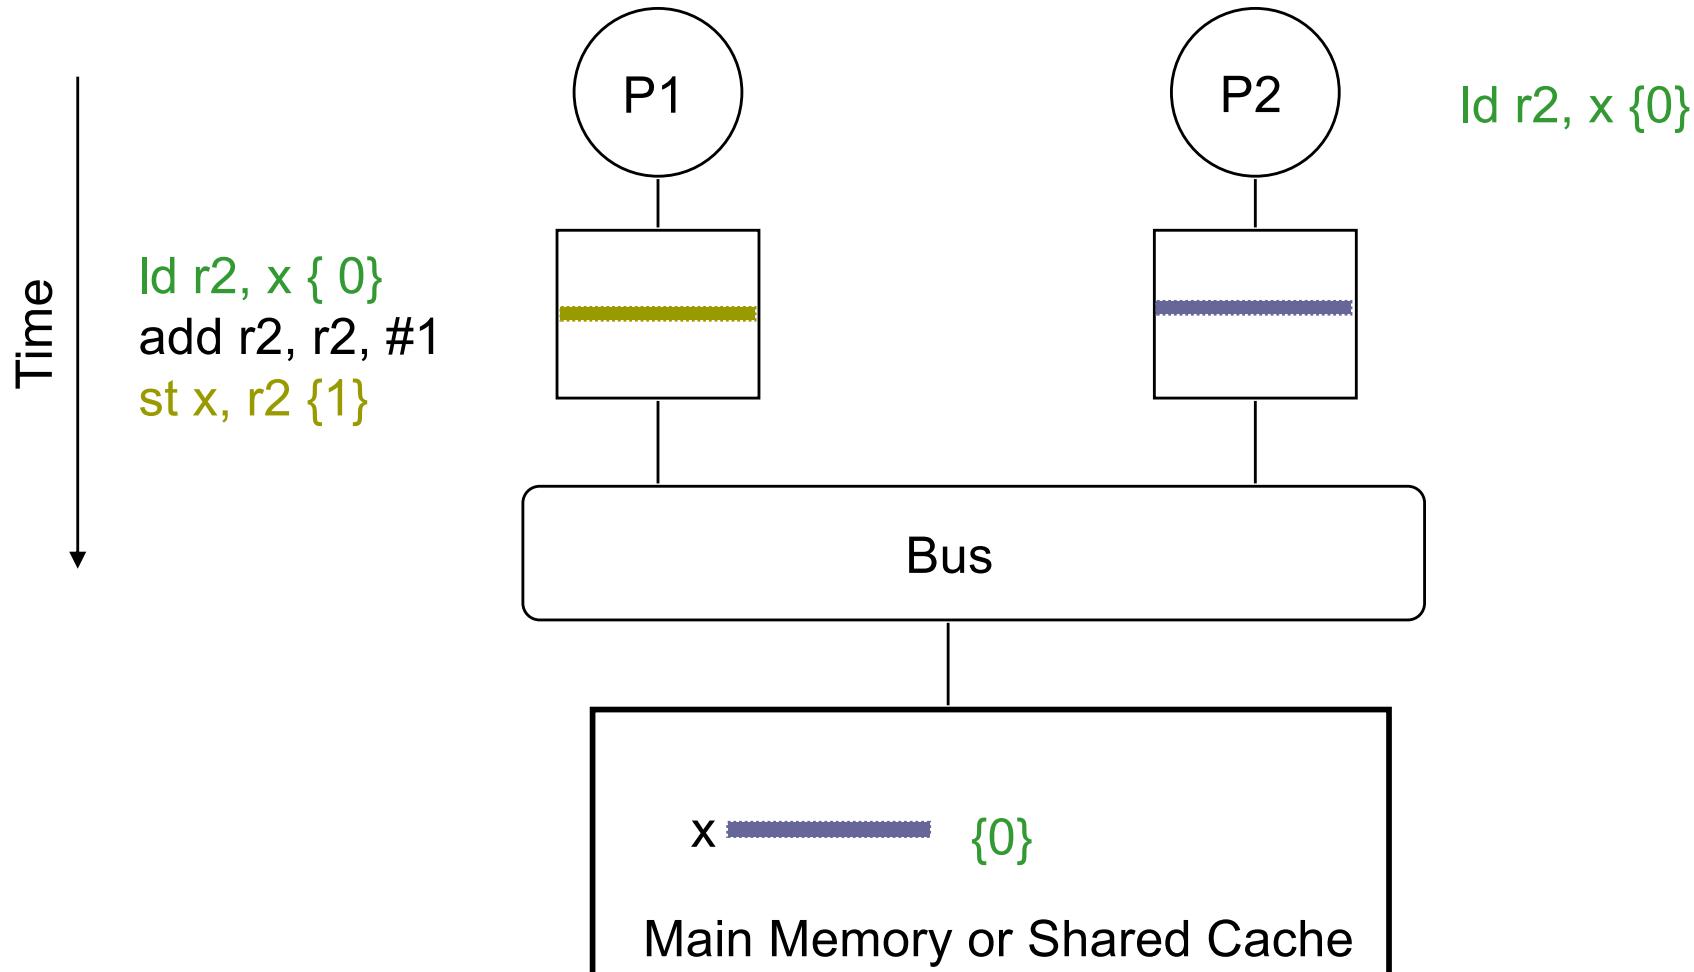
Cache Coherence Problem (Step 1)



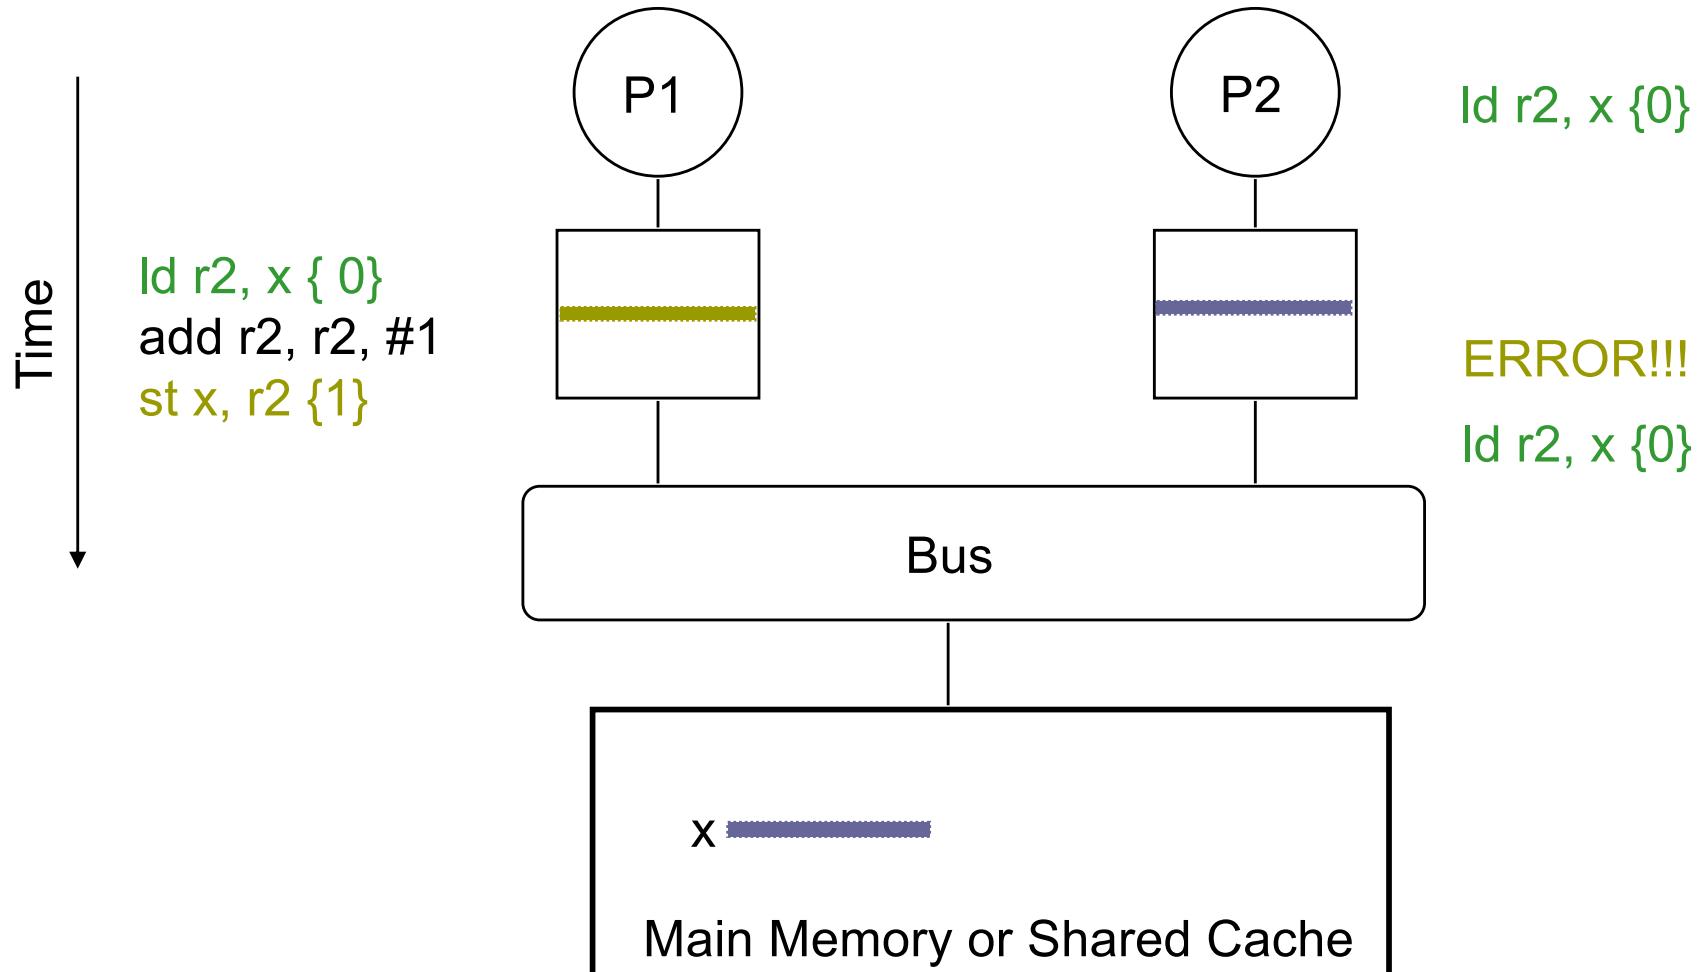
Cache Coherence Problem (Step 2)



Cache Coherence Problem (Step 3)



Cache Coherence Problem (Step 4)



Cache Coherence Protocols



- Cache coherence protocols
 - State and sequence of actions needed to ensure caches are coherence in multi-core systems
- Naïve protocol: turn off caching
 - Does not allow caching of read-shared data
 - Does not allow caching of private data
- Smarter protocol: allow one writer at the time
 - All caches can have copies of data while read-shared
 - On a write, invalidate all copies, allow a single writer



Cache Coherence Protocol

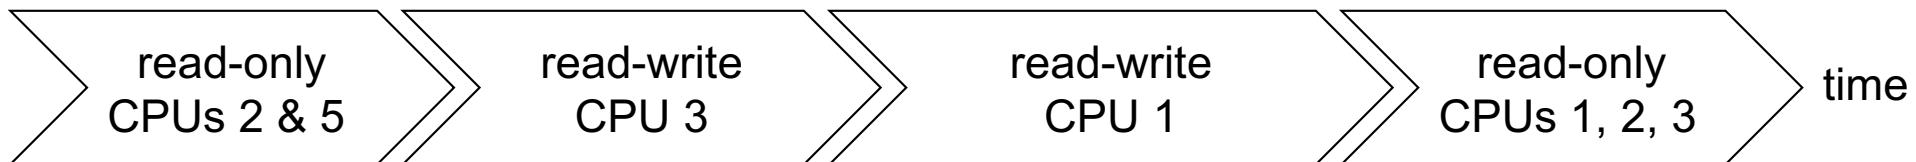
Divide lifetime of a memory value into epochs (time periods)

1. Single-Writer, Multiple-Read (SWMR) Invariant

For any memory location A, at any given time (epoch), there exists only a single CPU that may write to A (and can also read it) or some number of CPUs that may only read A

2. Data-Value Invariant

The value of the memory location at the start of an epoch is the same as the value of the memory location at the end of its last read-write epoch

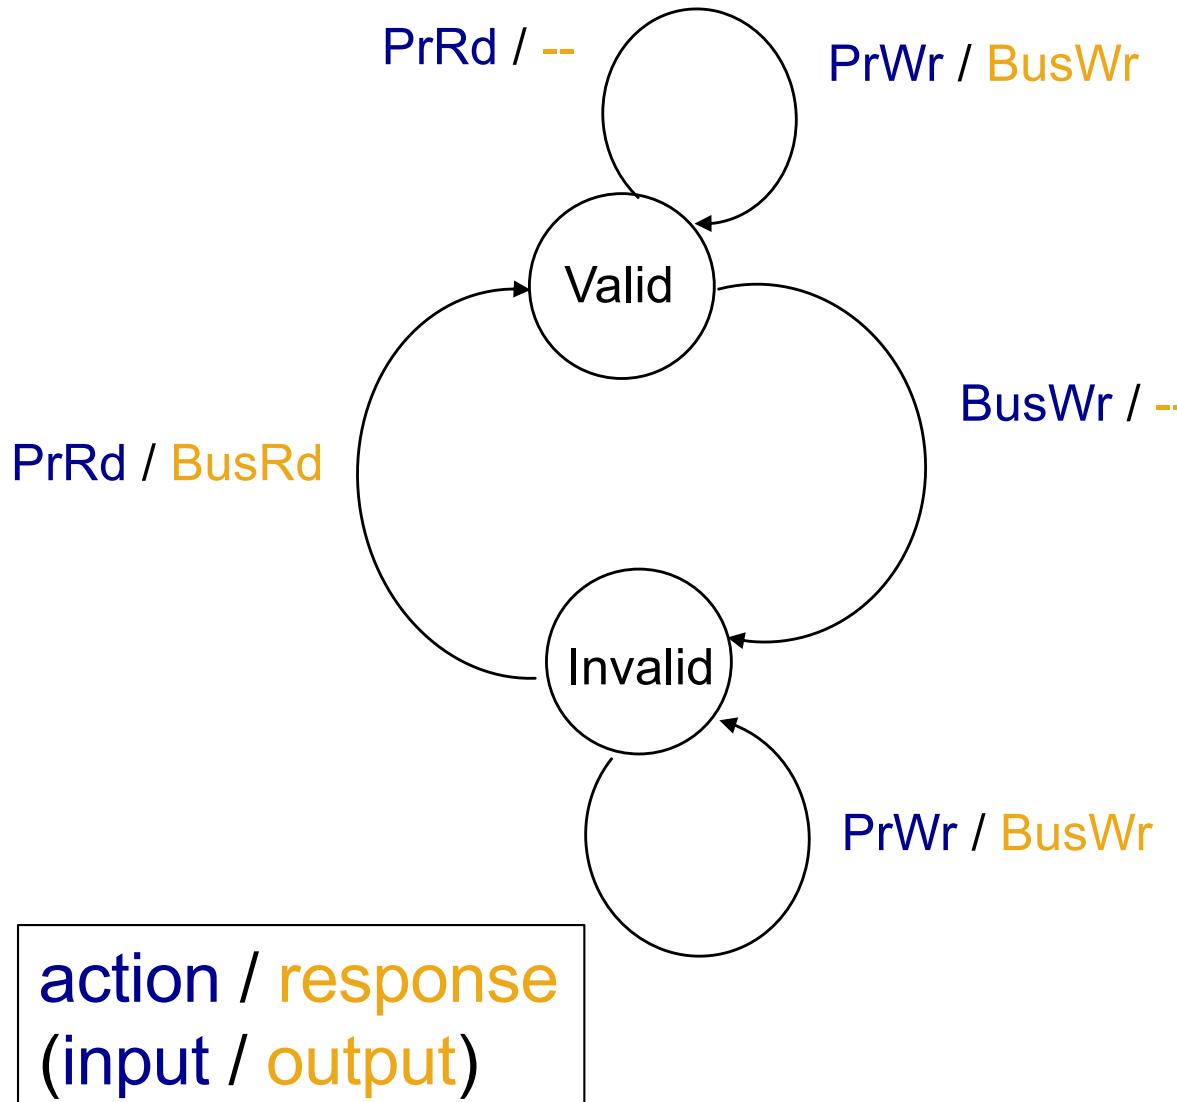


Implementation of Coherence Protocols



- Protocols rely on monitoring core actions
 - For now, assume all misses visible to all cores
 - Interconnect makes all actions visible
 - Cache FSM also reacts to requests from other caches
 - Known as a snooping-based protocols
- Protocols avoid stale copies
 - Whenever a core becomes a writer: invalidate copies

Simple Coherence Protocol



For write-through, no-write-allocate caches

Action	Abbreviation
Processor Read	PrRd
Processor Write	PrWr
Bus Read	BusRd
Bus Write	BusWr



Is 2-state Protocol Coherent?

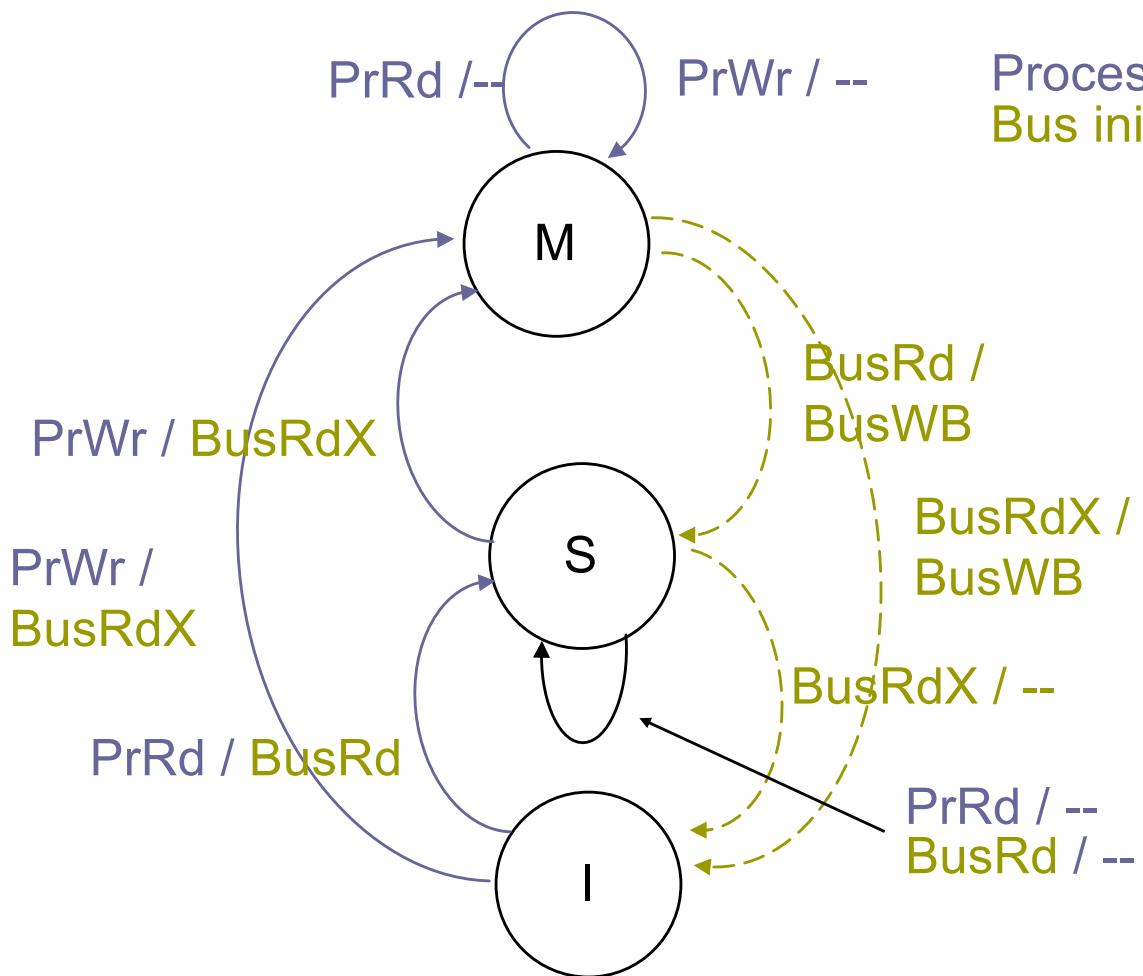
- Assume bus transactions and memory operations are atomic
 - Processor waits for memory operation to finish before issuing next
 - With one-level cache, assume invalidations applied during bus transaction
- SWMR Invariant
 - All writes go to bus + atomicity, causes transition to invalid state
- Data-Value Invariant
 - Read misses appear on bus, and will “see” last write in bus order
 - Read hits: do not appear on bus
 - But value read was placed in cache by either
 - Most recent write by this processor or most recent read miss
 - Both these transactions appeared on the bus
 - So reads hits also see values as produced bus order

A 3-State Write-Back Invalidation Protocol



- 2-State protocol
 - + Simple hardware and protocol (used in Sun Niagara)
 - - Bandwidth: write-through, every write goes on bus
 - Can work in multi-core with additional level of shared cache
- 3-State protocol (MSI)
 - Modified (called Exclusive in some books)
 - One cache has valid/latest copy
 - Memory is stale
 - Shared
 - One or more caches (and memory) have valid copy
 - Invalid
- Must invalidate all other copies before entering modified state
 - Requires bus transaction (order and invalidate)

MSI Coherence Protocol



Processor initiated
Bus initiated

Abbreviation	Action
PrRd	Processor Read
PrWr	Processor Write
BusRd	Bus Read
BusRdX	Bus Read Exclusive
BusWB	Bus Writeback