

# Scheduling Lecture

## Scheduling

- what is scheduling?
  - Goals
  - Mechanisms
- Scheduling on batch systems
- scheduling on interactive systems
- other kinds of scheduling
  - real-time scheduling

## Why Scheduling processes?

- Busts of CPU usage alternate with periods of I/O wait
- some processes are CPU-bound: they don't give many I/O requests
- other processes are I/O bound and make kernel requests

CPU bound processes: computing stuff

I/O bound processes: waits for input, sleeps a lot

1 second CPU time = 1 week

## When are processes scheduled?

- at the time they enter the system
  - common in batch systems
  - the types of batch scheduling:
    - submission of a new job causes the scheduler to run
    - scheduling only done when a job voluntarily gives up the CPU
- At relatively fixed intervals (clock interrupt)
  - necessary for interactive systems
  - may also be used for batch systems
  - scheduling algorithm at each interrupt, and picks the next process from the pile of "ready" processes

## Scheduling Goals:

- All systems:
  - Fairness: give each process a fair share of the CPU
  - Enforcement: ensure that the stated policy is carried out

- Balance: keep all parts of the system busy
- Batch systems:
  - throughput: maximize jobs per unit time (hour)
  - Turnaround time: minimize time users wait for jobs
  - CPU utilization: keep the CPU as busy as possible
- Interactive Systems:
  - response time: respond quickly to users' request
  - Proportionality: meet users' expectations
- Real-Time Systems:
  - meet deadlines: missing deadlines is a system failure
  - Predictability: same type of behavior for each time slice

Interactive vs batch scheduling:

- batch
  - First-Come-First-Served
  - Shortest Job First
  - Shortest Remaining Time First
  - Priority
- Interactive:
  - Round-Robin
  - Priority (preemptive)
  - Multi-Level Feedback Queue
  - Lottery Scheduling

First Come, First Served (FCFS):

- goal: do jobs in the order they arrive
  - fair in the same way a bank teller line is fair
- Problem: long bond delay every job after them

Shortest job first:

- goal: do the shortest job first
- optimal for throughput
  - throughput: number of jobs accomplished
- jobs sorted in increasing order of execution time
- shortest remaining time first: preemptive form of SJF:
  - re-evaluate when a new jobs is submitted
- problem: how does the scheduler know how long a job will take
- starvations: a job with a high execution time will keep getting pushed behind the line

as faster jobs are introduced

### Three-level Scheduling

- jobs held in input queue will be moved into memory
  - pick “complementary jobs” small and large, CPU and I/O intensive
  - jobs moved into memory when admitted
- CPU scheduler picks next job to run
- Memory scheduler picks some jobs from main memory and move them to disk if sufficient memory space
- When process is in memory then the process is able to run

### Round\_robin scheduling

- scheduling interactive process
  - give each process a fixed time slot (quantum)
  - rotate through “ready” processes
    - have to select through the processes
    - only one process per core can run
  - each process makes some progress
- What’s a good quantum?
  - too short: many process switches hurt efficiency
  - too long: poor response to interactive requests
  - typical length: 10- 100 ms
- strict rotation: round robin
- context switching: switching through processes, each one gets a quantum

### Priority Scheduling:

- Assign a priority to each process
  - ready process with highest priority allowed to run
  - running process may be interrupted after its quantum expires
- priorities may be assigned dynamically
  - reduced when a process uses CPU time
  - increased when a process waits for I/O
- often, processes grouped into multiple queues based on priority, and run round-robin per queue
- starvation: one high priority process is always ready, and another process doesn’t get to run

### Shortest process next

- run the process that will finish the soonest
  - in interactive systems, job completion time is unknown!
- guess at completion time based on previous runs
  - update estimate each time a job is run
  - estimate is a combination of previous estimate and most recent run time
- not often used because round robin with priority works so well!

### Lottery scheduling

- give processes “tickets” for CPU time
  - more tickets -> higher share of CPU
- each quantum, pick a ticket at random
  - if there are  $n$  tickets, pick a number from 1 to  $n$ 
    - pseudo-random number is ok if it's a good RNG
  - Process holding the ticket gets to run for a quantum
  - this can be implemented efficiently without “real” tickets
    - track range of tickets belonging to each process
- Over the long run, each process gets the CPU  $m/n$  of the time if the process has  $m$  of the  $n$  existing tickets
- tickets can be transferred
  - cooperating processes can exchange tickets
  - clients can transfer tickets to a server so it can have a higher priority
  - parent (shell) can transfer tickets to a child process
- don't keep track of ticket numbers
- doesn't matter where the scheduler starts, everyone will still get to run
- random number generated: subtract the lottery ticket numbers until 0 or negative number and whatever process that lands on choose that process to run
- no starvation with this scheduling

### Scheduling in BSD4

- quantum is 100ms : longest that's ok for interactive scheduling
- scheduler is based on multi-level feedback queues
  - priority is based on two things
    - resource requirements: blocked threads have higher priority when rescheduled
    - previous CPU usage: CPU ...

### Calculating priority in BSD4:

- thread priority is set by
  - $\text{priority} = \text{MIN} + [\text{estcpu}/4] + 2\text{xnice}$ 
    - values above MAX are set to MAX
    - MIN = 160, MAX = 223
    - nice is set by the user to manually over thread priority
    - cstime is an estimate of the number of “ready” processes in the CPU when the calculation is made
      - has a bit of “memory” so it doesn’t change too quickly
      - estcpu is update each clock tick
    - higher numbers indicate lower priority: threads with lowest priority values are scheduled first
- thread priority is set every 40ms
- scheduling is more complex for multiprocessors...

Scheduling user-level threads:

- kernel picks a process to run next
- if you have user-level threads, the kernel chooses the process to run
- Thread table:
  - non preemptive scheduling

Scheduling kernel level threads:

- kernel schedules each thread
  - no restrictions on ordering