

Virtual Memory (Long's Lecture)

Paging and page table

- virtual addresses mapped to physical addresses
 - unit of mapping is called a page
 - all addresses in the same virtual page are in the same physical page
 - page table entry contains translation for a single page
- table translates virtual page number to physical page number
 - note all virtual memory has a physical page
 - not every physical page need be used
- Overlay: piece of software called overlay manager: to get something from memory
- Make everything the same size: so everything can fit anywhere
 - cut programs into the same size pieces
 - lie to the CPU, CPU thinks programs are next to each other -> MMU does this lie
 - programs are virtually contiguous to each other
 - programs are physically spread out from each other

What's in a page table entry?

- each entry in the page table contains
 - valid bit: set if this logical page number has a corresponding physical frame in memory
 - if not valid, remainder of PTE is irrelevant
- page frame number: page in physical memory -> the top half of addresses
- referenced bit: set if data on the page has been accessed
- dirty (modified) bit: set if the data on the page has been modified
- protection information

Mapping logical addresses to physical addresses:

- split the address from CPU into two pieces
 - page number (p)
 - page offset (d)
- page number
 - index into page table
 - page table contains base address of page in physical memory
- page offset
 - added to base address to get actual physical memory address
- page size = 2^d bytes

Two-Level Page Tables

- problem: page tables can be too large
 - 2^{32} in 4KB pages \rightarrow 1 million PTEs
 - worst for 64-bit addressing
- solution: use multi-level page tables
 - page size in first page table is large
 - PTE marked invalid in first page table needs no 2nd level page table
- 1st level page table has pointers to 2nd level page table
- 2nd level page table has actual physical page numbers in it
- goal: not to jump around: locality is key

How long do memory accesses take?

- assume the following times:
 - TLM lookup time = a (often zero - overlapped in CPU)
 - Memory access time = m
- Hit ration (h) is percentage of time that a logical page number is found in the TLB
 - larger TLB usually means higher h
 - TLB structure can affect h as well
- effective access time (an average) is calculated as:
 - $EAT = (m+a)h + (m+m+a)(1-h)$
 - $EAT = a + (2-h)m$
- ...

Inverted page table

- reduce page table size further: keep one entry for each frame memory
 - alternative: merge tables for pages in memory and on disk
- PTE contains
 - virtual address pointing to this frame
 - info about the process that owns this page
- Search page table by
 - hashing the virtual page number and process ID
 - Starting at the entry corresponding to the hash result
 - search until either the entry is found or a limit is reached
- Page frame number is index of PTE
- Improve performance by using more advanced hashing algorithms

