

# Week 2

Monday, April 8, 2019 9:31 AM

## Unix

- System calls
  - Ex: `ls > foo` := `ls` contents pipes to `foo`
  - Ex: `ls 2> foo` := `ls` `stderr` contents pipes to `foo`
- File descriptors
  - When you open a file, you need to know everything like buffers and access info
  - Stored with array of pointers
  - We care about 0,1,2 which point to `stdin`, `stdout`, `stderr`
  - Redirect `stdin`
    - `F = open(" ", ...)`
    - `Close(0)`
    - `Dup(f)`
    - `Close(f)`
  - Pipe
    - `Int p[2]`
    - `Pipe(p)`
    - `p[1] := write`
    - `p[0] := read`
  - `Ls | wc`
    - Parent needs to declare pipe so children can both access
    - `Ls` child `close(1) dup(p[1]) close(p[0])`
    - `Wc` child `close(0) dup(p[0]) close(p[1])`
    - Parent `close(0) dup(p[0])`

## What is a process

- Code, data, stack
  - Usually but not always has its own address space
- Program state
  - CPU registers
  - Program counter
  - Stack pointer
- Only one process can be running in a single cpu core at any given time
  - Multi-core CPUs can support multiple processes

## The process model

- Multiprogramming of four programs
- 4 independent processes
- Processes run sequentially
- Only one program active at any instant
  - Each instant can be incredibly short
  - Only applies if there is a single CPU with a single core in the system
- Context switching

## When is a process end

- Either voluntary or involuntary
  - Voluntary
    - `Exit` or `Error exit`
  - Involuntary
    - Fatal error
    - Killed by another process

## Process hierarchies

- Parents create child processes
  - o Childs can create their own child processes
- Unix calls these groups Process Groups
- If a Process is terminated, its children are inherited by the parent

## Process States

- Process is in one of 5 states
  - o Created
  - o Ready
  - o Running
  - o Blocked/waiting
  - o Exit
- Transition
  - o Process enters ready queue
  - o Scheduler picks this process
  - o Scheduler picks a different process
  - o Process waits for an event such as IO
  - o Event occurs
  - o Process Exits
  - o Process ended by another process

## Processes in the OS

- Two layers for processes
- Lowest layer of process-structured OS handles interrupts, scheduling
- Above that layer are sequential processes
  - o Processes tracked in the process table
  - o Each process has a process table entry

## Process table entry

- Process management
  - o Registers
  - o Program counter
  - o CPU status word
  - o Stack pointer
  - o Process state
  - o Priority
  - o Process id
  - o Parent id
  - o Signals
  - o Process start time
  - o Total CPU usage
- File management
  - o Root directory
  - o Current directory
  - o File descriptions
  - o User id
  - o Group id
- Memory management
  - o Pointers to text, data, stack
  - o Pointer to page table

## What happens on a trap/interrupt

- Hardware saves program counter

- Hardware loads new PC, identifies interrupt
- Assembly language routine saves registers
- Assembly language sets up the stack
- Assembly language calls C to run service routine
- Service routine calls scheduler
- Scheduler selects a process to run next might be interrupted
- Assembly language routine loads PC and registers for the selected process

Threads: "processes" sharing memory

- Process :: address space
- Thread :: program counter / stream of instructions
- Two examples
  - o Three processes each with one thread
  - o One process with three threads

Each thread needs its own stack

Why use threads?

- Allow a single application to do many things at once
  - o Simple programming model
  - o Less waiting
- Threads are faster to create or destroy
  - o No separate address space
- Overlap computation and IO
  - o Could be done without threads but its harder
- EX: word processor
  - o Thread to read from keyboard
  - o Thread to format document
  - o Thread to write to disk

Issue with threads

- May be tricky to convert single thread code to multithread code
- Re-entrant code
  - o Code must function properly when multiple threads are using it at same time
  - o Be careful of using static or global variables

User level threads

- + No need for kernel support
- - Maybe slower than kernel threads
- - Harder to do non-blocking IO

Kernel level threads

- + More flexible scheduling
- + Non-blocking IO
- - Not necessarily portable

Posix threads

- Standard interface for threading library