

# Monday

Monday, April 15, 2019 12:03 PM

## Bounded buffer

- Everything size n
- Producer
  - o Down empty
  - o Down mutex
  - o Input item
  - o Increase in by 1
  - o Up mutex
  - o Up full
- Consumer
  - o Down full
  - o Down mutex
  - o Set item to out
  - o Increase out
  - o Up mutex
  - o Up empty
- If, by any chance, empty and mutex lines switch
  - o The code will block up
  - o This might not always break
- Don't use semaphores
- Use monitors instead

Review SYNCHRONIZATION CODE  
Review BATCH COMPUTING

## Dining Philosophers problem

- N philosophers around a table
  - o All hungry and like to think
- N chopsticks available
  - o 1 between each pair of philosophers
- Need 2 chopsticks to eat
- Alternate between eating and thinking
- Goal: coordinate use of chopsticks
- Example of Deadlock
  - o Aborting a process by eliminating a "philosopher"
- Hold and wait
- NO preemption
- Circular wait
- Mutual exclusion

## Scheduling

- What is Scheduling
  - o Doing things in a certain order
  - o Which process do I run next
  - o Without loss of generality
  - o Goals/Mechanism

## Processes

- I/O bound vs CPU bound
- 1 second of cpu time is roughly 1 week of i/o bound
- Bursts of CPU usage alternate with periods of I/O wait
- Some processes are CPU-bound: they don't make many I/O requests
- Other processes are I/O-bound and make many kernel requests

## When are processes scheduled

- At the time they enter the system
  - o Common in batch systems
  - o Two types of batch scheduling
    - Submission of a new job causes the scheduler to run
    - Scheduling only done when a job voluntarily gives up the CPU
- At relatively fixed intervals

- Necessary for interactive systems
- May also be used for batch systems
- Scheduling algorithms at each interrupt, and picks the next process from the pool of "ready" processes

#### Scheduling goals

- All systems
  - Fairness
  - Enforcement
  - Balance
- Batch systems
  - Throughput
  - Turnaround time
  - CPU utilization
- Interactive systems
  - Response time
  - Proportionality
- Real-time systems
  - Meet deadlines
  - Predictability

#### Interactive vs Batch

- Batch
  - First Come First Served (FCFS)
  - Shortest Job First (SJF)
  - Shortest Remaining Time First (SRTF)
  - Priority (non-preemptive)
- Interactive
  - Round-Robin (RR)
  - Priority (preemptive)
  - Multi-level feedback queue
  - Lottery scheduling

# Wednesday

Wednesday, April 17, 2019 11:50 AM

## FCFS

- Do jobs in the order they arrive
- Very simple algorithm
- Problem
  - o Long jobs delay every job after them
  - o Many jobs may be delayed for a long time

## SJF

- Do shortest jobs first
  - o Short jobs first
  - o Long jobs delay jobs after them
- Jobs sorted in increasing order of execution time
- SRTF: preemptive form of SJF
  - o Re-evaluate when a new job is submitted
- Problem:
  - o How does the scheduler know how long a job will take
  - o Starvation
    - When shorter jobs are introduced, the longer jobs keep getting delayed

## Three level scheduling

- Jobs held in input queue until moved into memory
  - o Pick complementary jobs: small and large, CPU and I/O intensive
  - o Jobs move into memory when admitted
- CPU scheduler picks next job to run
- Memory scheduler picks some jobs from main memory and moves them to disk if insufficient memory space

## Round Robin

- Scheduling interactive processes
  - o Give each process a fixed time slot
    - Quantum
  - o Rotate through ready processes
  - o Each process makes some progress
- What's a good quantum
  - o Too short
    - Many process switches hurt efficiency
  - o Too long
    - Poor response to interactive requests
  - o Typical length = 10 to 100 ms
- Strict rotation

## Priority Scheduling

- Assign a priority to each process
  - o Ready process with the highest priority allowed to run
  - o Running process may be interrupted after its quantum expires
- Priorities may be assigned dynamically

- Reduced when a process uses CPU time
  - Increased when a process waits for IO
- Often processes grouped into multiple queues based on priority, and run round-robin per queue
- When quantum runs out, we punish the process like decreasing its priority
  - Policy = how we said things should happen
  - Mechanism - enforcement

#### Shortest process next

- Run the process that will finish the soonest
  - In interactive systems, job completion time is unknown
- Guess at completion time based on previous runs

#### Lottery Scheduling

- Give processes tickets for CPU time
  - More tickets = higher share of CPU
- Each quantum, pick a ticket at random
  - If there are n tickets, pick a number from 1 to n
    - Pseudo-random number is OK if it's a good RNG
  - Process holding the ticket gets to run for a quantum
  - This can be implemented efficiently without real tickets
    - Track range of tickets belonging to each process
- Over the long run, each process gets the CPU m/n of the time if the process has m of the n existing tickets
- Tickets can be transferred
  - Cooperating processes can exchange tickets
  - Clients can transfer tickets to a server so it can have a higher priority
  - Parent (shell) can transfer tickets to a child process

#### Scheduling in BSD4

- Quantum is 100 ms:
  - Longest that's OK for interactive scheduling
- Scheduler is based on multilevel feedback queues
  - Priorit is based on two things
    - Resource requirements

#### Calculate priority

- Thread priority is set by:
  - $Pri = MIN + [estcpu/4] + 2 * nice$
  - Values above MAX are set to MAX
  - MIN = 160 MAX = 223
  - Nice is set by the user to manually lower thread priority
  - Estcpu is an estimate of the number of ready processes in the CPU when the calculation is made
    - Has a bit of memory so it doesn't change too quickly
    - Estcpu is updated each clock tick
  - Higher numbers indicate lower priority: threads with lowest priority values are scheduled first
- Thread priority is set every 40 ms
- Scheduling is more complex for multiprocessors