

Part1

Monday, May 6, 2019

12:02 PM

Least Recently Used

- Assume pages used recently will be used again soon
 - o Throw out page that has been unused for longest time
- Must keep a linked list of pages
 - o Most recently used at front, least at rear
 - o Update this list every memory reference!
 - This can be somewhat slow: hardware has to update a linked list on every reference!
- Alternatively, keep counter in each page table entry
 - o Global counter increments with each CPU cycle
 - o Copy global counter to PTE counter on a reference to the page
 - o For replacement, evict page with lowest counter value
- LRU stack, put most recently used onto top of "stack", look through it and put last used to top of stack

Simulating LRU in software

- Few computers have the necessary hardware to implement full LRU
 - o Linked list method impractical in hardware
 - o Counter-based method could be done, but it's slow to find the desired page
- Approximate LRU with Not Frequently used NFU algorithm
 - o At each clock interrupt, scan through page table
 - o If $R=1$ for a page, add one to its counter value
 - o On replacement pick the page with the lowest counter value

Aging replacement algorithm

- Reduce counter values over time
 - o Divide by two every clock cycle, use right shift
 - o More weight given to more recent references
- Select page to be evicted by finding the lowest counter value
- Algorithm is:
 - o Every clock tick, shift all counters right by 1 bit
 - o On reference, set leftmost bit of a counter, can be done by copying the reference bit to the counter at the clock tick

Working set

- Demand paging: bring a page into memory when it's requested by the process
- How many pages are needed?
 - o Could be all of them, but not likely
 - o Instead, processes reference a small set of pages at any given time-locality of reference
 - o Set of pages can be different for different processes or even different times in the running of a single process
- Set of pages used by a process in a given interval of time is called the working set
 - o If entire working set is in memory, no page faults!
 - o If insufficient space for working set, thrashing may occur
 - o Goal: keep most of working set in memory to minimize the number of page faults suffered by a process

How big is the working set?

- Working set is the set of pages used by the k most recent memory references
- $W(k,t)$ is the size of the working set at time t
- Working set may change over time
 - o Size of working set can change over time as well...

Modeling page replacement algorithms

- Goal: provide quantitative analysis or simulation showing which algorithms do better
 - o Workload page reference string is important: different strings may favor different algorithms
 - o Compare algorithms to one another
 - o Model parameters within an algorithm
 - Number of available physical pages
 - Number of bits for aging

How is modeling done?

- Generate a list of references
 - o Artificial (made up)
 - o Trace a real workload (set of processes)
- Use an array (or other structure) to track the pages in physical memory at any given time
 - o May keep other information per page to help simulate the algorithm (modification time, time when paged in, etc.)
- Run through references, applying the replacement algorithm

Belady's anomaly

- Reduce the number of page faults by supplying more memory
 - o Use previous reference string and FIFO algorithm
 - o Add another page to physical memory (total 4 pages)
- More page faults, not fewer!
 - o Adding more pages shouldn't result in worse performance
- Motivated the study of paging algorithms

Modeling more replacement algorithm

- Paging system characterized by:
 - o Reference string of executing process
 - o Page replacement algorithm
 - o Number of page frames available in physical memory (m)
- Model this by keeping track of all n pages references in array M
 - o Top part of M has m pages in memory
 - o Bottom part of M has $n-m$ pages stored in disk
- Page replacement occurs when page moves from top to bottom

LRU example

- Model LRU replacement with
 - o 8 unique references in the reference string
 - o 4 pages of physical memory
- Array state over time shown below
- LRU treats list of pages like a stack

Stack Algorithms

- LRU is an example of a stack algorithm
- For stack algorithms
 - o Any page in memory with m physical pages is also in memory with $m+1$ physical pages
 - o Increasing memory size is guaranteed to reduce (or at least not increase) the number of

page faults

- Stack algorithms do not suffer from Belady's anomaly
- Distance of a reference \Leftrightarrow position of the page in the stack before the reference was made
 - o Distance is infinity if no reference had been made before
 - o Distance depends on reference string and paging algorithm: might be different for LRU and optimal (both stack algorithms)

Predicting page fault rates using distance

- Distance can be used to predict page fault rates
- Make a single pass over the reference string to generate the distance string on the fly
- Keep an array of counts
 - o Entry j counts the number of times distance j occurs in the distance string
- The number of page faults for a memory of size m is the sum of counts for $j > m$
 - o This can be done in a single pass
 - o Makes for fast simulations of page replacement algorithms
- This is why virtual memory theorists like stack algorithms

Local vs. global allocation policies

- What is the pool of pages eligible to be replaced
 - o Pages belonging to the process needing a new page
 - o All pages in the system
- Local allocation: replace a page from this process
 - o May be more fair: penalize processes that replace many pages
 - o Can lead to poor performance: some processes need more pages than others
- Global allocation: replace a page from any process

Control overall page fault rate \Rightarrow **THRASHING**

- Despite good designs, system may still thrash
- Most (or all) processes have high page fault rate
 - o Some processes need more memory, ...
 - o But no processes need less memory (and could give some up)
- Problem: no way to reduce page fault rate
- Solution: Reduce number of processes competing for memory
 - o Swap one of more to disk, divide up pages they held
 - o Reconsider degree of multiprogramming

Part 2

Wednesday, May 8, 2019 12:20 PM

How big should a page be?

- Smaller pages have advantages
 - o Less internal fragmentation
 - o Better fit for various data structures, code sections
 - o Less unused physical memory (some pages have 20 useful bytes and the rest isn't needed currently)
- Larger pages are better because
 - o Less overhead to keep track of them
 - Smaller page tables
 - TLB can point to more memory (same number of pages, but more memory per page)
 - Faster paging algorithms (fewer table entries to look through)
 - o More efficient to transfer larger pages to and from disk

Separate I and D address spaces

- One user address space for both data and code
 - o Simpler
 - o Code/data separation harder to enforce
 - o More address space?
- One address space for data, another for code
 - o Code and data separated
 - o More complex in hardware
 - o Less flexible
 - o CPU must handle instructions and data differently
- FreeBSD does the former

Sharing pages

- Processes can share pages
 - o Entries in page table point to the same physical page frame
 - o Easier to do with code: no problems with modification
- Virtual addresses in different processes can be
 - o The same: easier to exchange pointers, keep data structures consistent
 - o Different: may be easier to actually implement
 - Not a problem if there are only a few shared regions
 - Can be very difficult if many processes share regions with each other

Shared libraries

- Many libraries are used by multiple programs
- Only want to keep a single copy in memory
- Two possible approaches
 - o Fixed address in memory
 - No need for code to be relocatable
 - How can libraries be placed?
 - o Per-process address in memory
 - More flexible: no central arbiter of addresses
 - Code has to be relocatable...

Memory mapped files

- Extension of shared libraries
 - o File blocks mapped directly into a process's address space
 - o Process can access file data just like memory
- Advantages
 - o Efficient: no need for read() and write() calls
 - OS manages 'paging' blocks in and out
 - o Easy to program: no buffer management
 - Can handle very large files, too
- Disadvantages
 - o Added burden for the OS: may not manage as well as program could
 - o Difficult to specify order in which writes are flushed to disk
 - OS writes pages back in unpredictable order
 - o Shared files can cause issues
 - Might be mapped to different addresses in different processes
 - Write ordering matter even more!
- Pointers don't make sense in this shared space

When are dirty pages written to disk

- On demand (when they're replaced)
 - o Fewest writes to disk
 - o Slower: replacement takes twice as long (must wait for disk write and disk read)
- Periodically (in the background)
 - o Background process scans through page tables, writes out dirty pages that are pretty old
- Background process also keeps a list of pages ready for replacement
 - o Page faults handled faster: no need to find space on demand
 - o Cleaner may use the same structures discussed earlier (clock, etc.)

Implementation issues

- Four times when OS involved with paging
- Process creation
 - o Determine program size
 - o Create page table
- During process execution
 - o Reset the MMU for anew process
 - o Flush the TLB (or reload it from saved state)
- Page fault time
 - o Determine virtual address causing fault
 - o Swap target page out, needed page in
- Process termination time
 - o Release page table
 - o Return pages to the free pool

How is a page fault handled?

- Hardware causes a page fault
- General registers saved (as on every exception)
- OS determines which virtual page needed
 - o Actual fault address in a special register
 - o Address of fualting instruction in register
 - Page fault was in fetching instruction, or
 - Page

Part 3

Friday, May 10, 2019 12:03 PM

Storing pages on disk

- Pages removed from memory are store on disk
- Where are they placed?
 - o Static swap area: easier tp code, less flexible
 - o Dynamically allocated space: more flexible, harder to locate a page
 - Dynamic placement often uses a special file (managed by the file system) to hold pages
- Need to keep track of which pages are where within the on-disk storage

Separating policy and mechanism

- Mechanism for page replacement has to be in kernel
 - o Modifying page tables
 - o Reading and writing page table entries
- Policy for deciding which pages to replace could be in user space
 - o More flexibility

Segmentation

- Different units in a single virtual address space
 - o Each unit can grow
 - o How can they be kept apart?
- Solution: segmentation
 - o Give each unit its own address space
- Segmentation registers

Using segments

- Each region of the process has its own segment
- Each segment can start at 0
 - o Addresses within the segment relative to the segment start
- Virtual addresses are <segment #, offset within segment>

Memory management in x86-32

- Memory composed of segments
 - o Segment pointed to by segment descriptor
 - o Segment selector used by identify descriptor
- Segment descriptor describes segment
 - o Base virtual address
 - o Size
 - o Protection
 - o Code / data