

Part1

Monday, May 13, 2019 12:00 PM

Device Controllers

- I/O devices have components
 - o Mechanical components
 - o Electronic components
- Electronic component controls the device
 - o May be able to handle multiple devices
 - o May be more than one controller per mechanical component
 - o Such as hard drive
- Controller's tasks
 - o Convert serial bit stream to block of bytes
 - o Perform error correction as necessary
 - o Make available to main memory

How is memory mapped I/O done?

- Single bus
 - o All memory accesses go over a shared bus
 - o I/O and RAM accesses compete for bandwidth
- Dual Bus
 - o RAM access over high speed bus
 - o I/O access over low speed bus
 - o Less competition for bus
 - o More hardware (more expensive)

I/O software: goals

- Device independence
 - o Programs can access any I/O device
 - o No need to specify device in advance
- Uniform naming
 - o Name of a file or device is a string or an integer
 - o Doesn't depend on the machine (underlying hardware)
- Error handling
 - o Done as close to the hardware as possible
 - o Isolate higher-level software
- Synchronous vs asynchronous transfers
 - o Blocked transfers vs interrupt-driven
- Buffering
 - o Data coming off a device cannot be stored in final destination
- Sharable vs dedicated devices

Interrupt handlers

- Interrupt handlers are best hidden
 - o Driver starts an I/O operation and blocks
 - o Interrupt notifies of completion
- Interrupt procedure does its task
 - o Then unblocks driver that started it
 - o Perform minimal actions at interrupt time
 - Some of the functionality can be done by the driver after it is unblocked

- Interrupt handler must
 - o Save registers not already save by interrupt hardware
 - o Set up context for interrupt service procedure

Device drivers

- Device drivers go between device controllers and rest of OS
 - o Drivers standardize interface to widely varied devices
- Device drivers communicate with controllers over bus
 - o Controllers communicate with devices themselves

Disk drive structure

- Data stored on surfaces
 - o Up to two surfaces per platter
 - o One or more platters per disk
- Data in concentric tracks
 - o Tracks broken into sectors
 - 256B - 4KB per sector
 - o Cylinder: corresponding tracks on all surfaces
- Data read and written by heads
 - o Actuator moves heads
 - o Heads move in unison

Part2

Friday, May 17, 2019 12:04 PM

Disk Request scheduling

- Goal: use disk hardware efficiently
 - o Bandwidth as high as possible
 - o Disk transferring as often as possible (and not seeking)
- We want to
 - o Minimize disk seek time (moving from track to track)
 - o Minimize rotational latency (waiting for disk to rotate the desired sector under the read/write head)
- Calculate disk bandwidth by
 - o Total bytes transferred / time to service request
 - o Seek time and rotational latency are overhead (no data is transferred), and reduce disk bandwidth

Disk scheduling algorithms

- Schedule disk requests to minimize disk seek time
 - o Seek time increases as distance increases (though not linearly)

First Come First serve

- Requests serviced in the order in which they arrived
 - o Easy to implement
 - o May involve lots of unnecessary seek distance
- Seek order
 - o 100 175 51 133 8 140 73 77
 - o Start 63
 - o Total cylinders = 646

Shortest seek time first

- Service the request with the shortest seek time from the current head position
 - o Form of SJF scheduling
 - o May starve some requests
- Seek order
 - o 73 77 51 8 100 133 140 175
 - o Start 63
 - o Total cylinders = 250

SCAN (elevator algorithm)

- Disk arm starts at one end of the disk and moves towards the other end servicing requests as it goes
 - o Reverses direction when it gets to end of the disk
 - o Also known as elevator algorithm
- Seek order
 - o 51 8 0 73 77 100 133 140 175
 - o Start 63
 - o Total cylinders = 238

C-SCAN

- Identical to SCAN except head returns to cylinder 0 when it reaches the end of the disk

- Treats cylinder list as a circular list that wraps around the disk
 - Waiting time is more uniform for cylinders near the edge of the disk
- Seek order
 - 73 77 100 133 140 175 199 0 8 51
 - Start 63
 - Total cylinders = 386

C-LOOK

- Identical to C-SCAN except head only travels as far as the last request in each direction
 - Saves seek time from last sector to end of disk
- Seek order
 - 73 77 100 133 140 175 8 51
 - Start 63
 - Total cylinders = 322

Picking a disk scheduling algorithm

- SSTF is easy to implement and works OK if there aren't too many disk requests in the queue
- SCAN-type algorithms perform better for systems under heavy load
 - More fair than SSTF
 - Use LOOK rather than SCAN algorithms to save time
- Long seeks aren't too expensive, so choose C-LOOK over LOOK to make response time more even
- Disk request scheduling interacts with algorithms for allocating blocks to files
 - Make scheduling algorithm modular: allow it to be changed without changing the file system
- Use SSTF or FCFS for lightly loaded systems
- Use C-LOOK for heavily loaded systems

Flash memory and SSDs

- Compared to disk, flash is
 - Faster (shorter access time, but lower bandwidth)
 - More expensive
 - More reliable (devices)
 - Less reliable (sectors)
- Compared to DRAM, flash is
 - Cheaper (a bit)
 - Non-volatile (data survives power loss)
 - Slower
- Use flash as a level between disk and memory?
- Issues
 - Can't overwrite sectors "in place"
 - Flash wears out: can only write 3,000 to 100,000 times per memory cell

Writing to flash memory

- Can't overwrite page in place
 - Need to write to "clean" region
 - Unit of erase is "erase block"
- One solution: copy entire erase block to new location
 - Modify the page being written
- This isn't efficient!
 - Flash vendors spend a lot of time on Flash Translation layers
 - FTLs map logical page numbers to physical pages
 - Also manage wear leveling: distribute erases among blocks

FTL lies. It makes it look like a disk

Is there a better way to manage flash

- Log-structured data storage works well on flash
 - o Append updates to a "log"
- Periodically "clean" structures
 - o Rewrite live data to a new log segment
 - o Obsolete data isn't rewritten
- Need to track mapping of logical locks to physical locations

Clock hardware

- Crystal oscillator with fixed frequency (only when computer is on)
 - o Typically used too time short intervals (~1 second)
 - o May be used to correct time-of-day clock
- Time-of-day clock (runs when system is off)
 - o Keeps minutes, hours, days
 - o May not be too accurate...
 - o Used to load system clock at startup
- Time kept in seconds and ticks (often 10^2 to 10^6 per second)
 - o Number of seconds since a particular time
 - For many versions of Unix, tick 0 was on January 1, 1970
 - o Number of ticks this second
 - o Advance ticks once per clock interrupt
 - o Advance seconds when ticks "overflow"

Use timers in software

-